

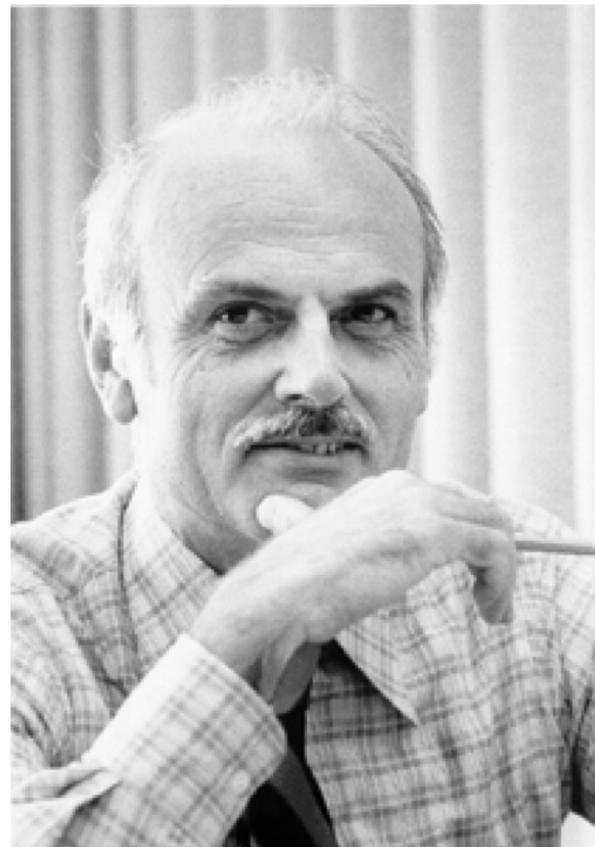
Ghislain Fourny

Big Data for Engineers Spring 2020

2. Lessons learnt from the past



Mr. Databases: Edgar Codd

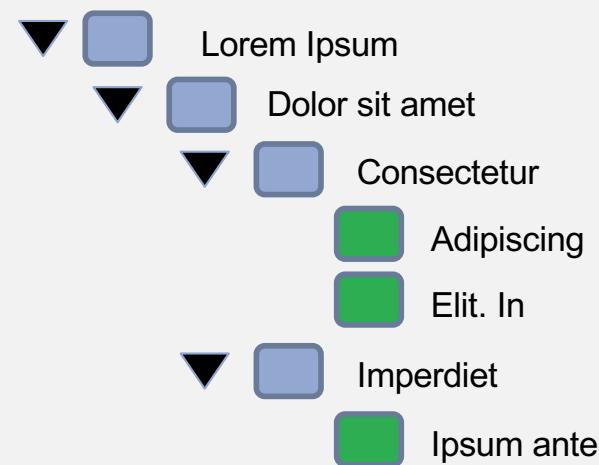


Wikipedia

Data Independence (Edgar Codd)

Logical data model

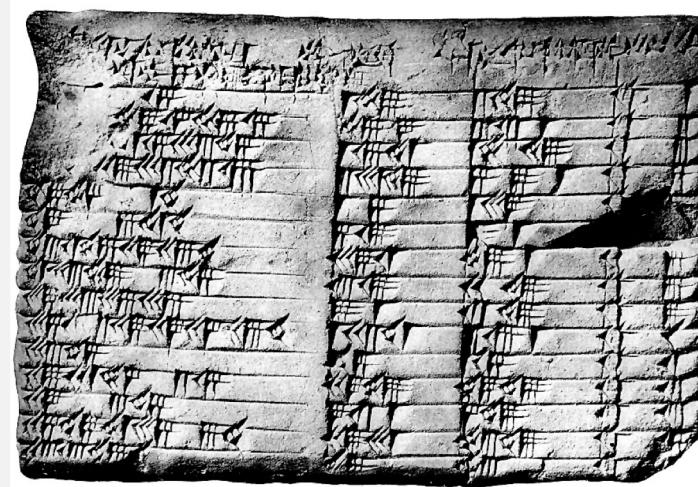
Physical storage



Data Independence (Edgar Codd)

Logical data model

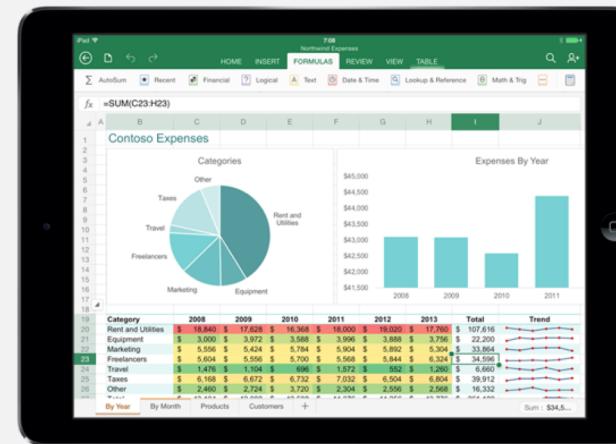
Physical storage



Data Independence (Edgar Codd)

Logical data model

Physical storage



Data Independence (Edgar Codd)

Logical data model

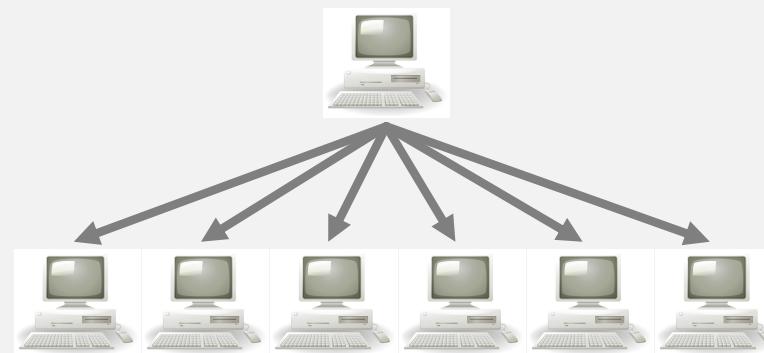
Physical storage



Data Independence (Edgar Codd)

Logical data model

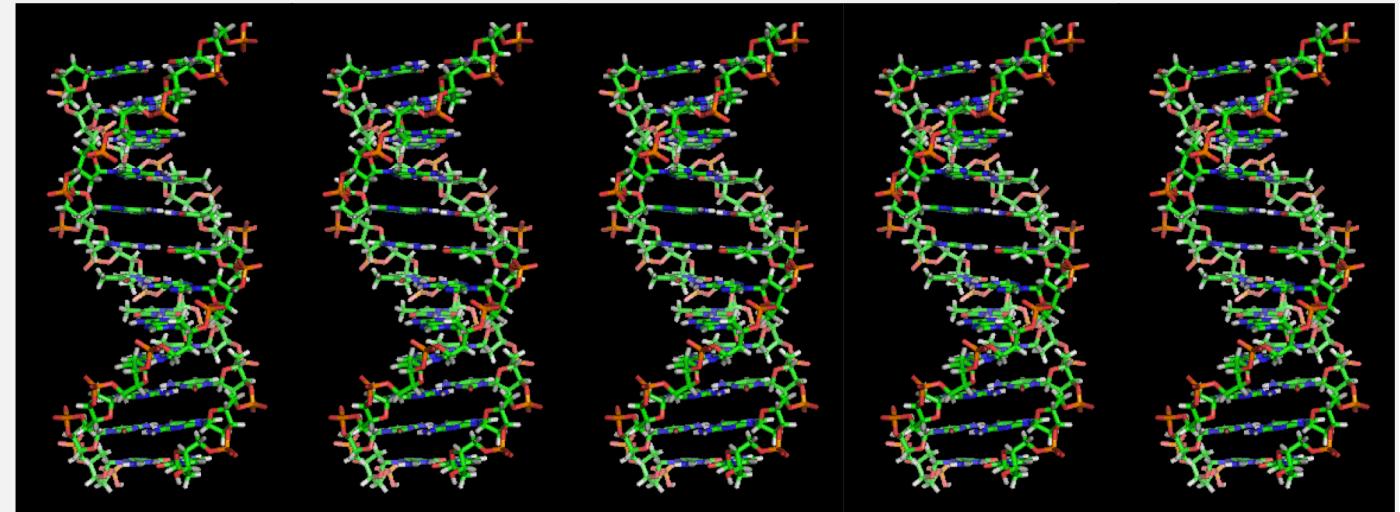
Physical storage



Data Independence (Edgar Codd)

Logical data model

Physical storage

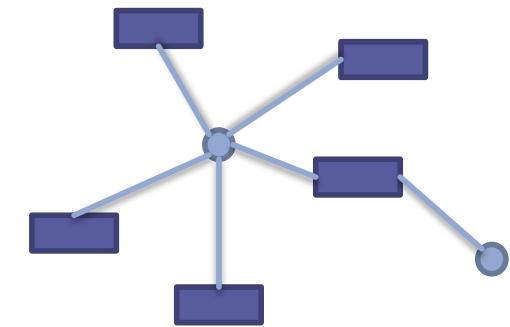
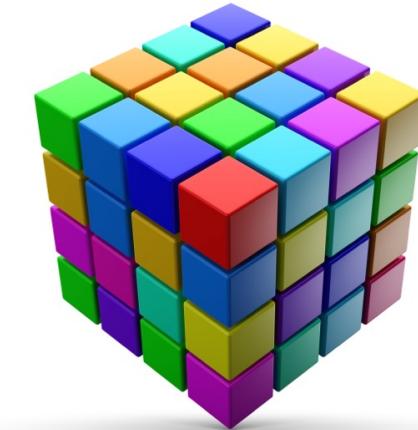


Data Shapes

Etiam vel erat nec dui aliquet vulputate sed quis nulla. Donec eget ultricies magna, eu dignissim elit. Nullam sed urna nec nisl rhoncus ullamcorper placerat et enim. Integer varius ornare libero quis consequat. Etiam a purus suscipit, accumsan nibh vel, posuere ipsum. Nulla nec tempor nibh, id venenatis lectus. Duis lobortis id urna eget tincidunt.

Praesent nec libero metus. Praesent at turpis placerat, congue ipsum eget, scelerisque justo. Ut volutpat, massa ac lacinia cursus, nisl dui volutpat arcu, quis interdum sapien turpis in tellus. Suspendisse potenti. Vestibulum pharetra justo massa, ac venenatis mi condimentum nec. Proin viverra tortor non orci suscipit rutrum. Phasellus sit amet euismod diam. Nullam convallis nunc sit amet diam suscipit dapibus. Integer porta hendrerit nunc. Quisque pharetra congue porta. Suspendisse vestibulum sed mi in euismod. Etiam a purus suscipit, accumsan nibh vel, posuere ipsum. Nulla nec tempor nibh, id venenatis lectus. Duis lobortis id urna eget tincidunt.

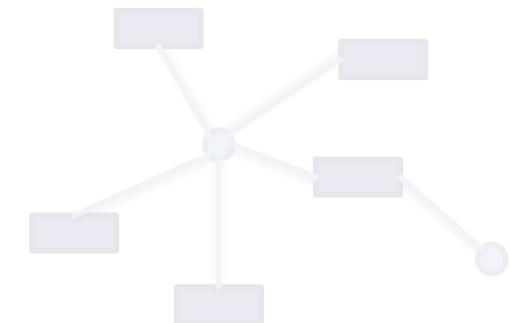
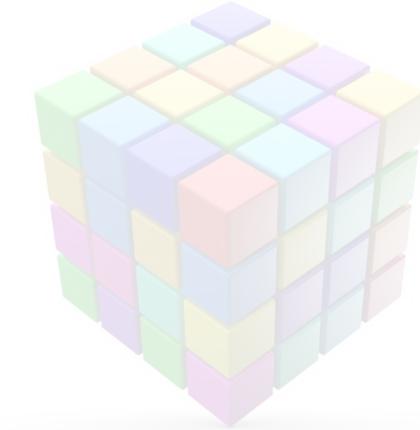
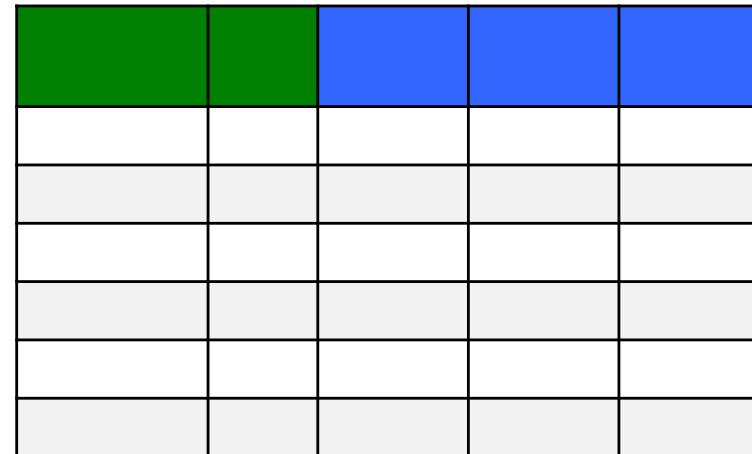
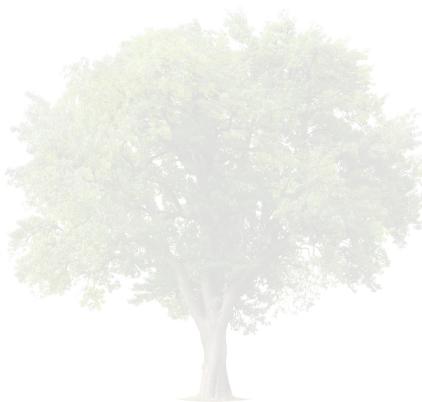




Data Shapes

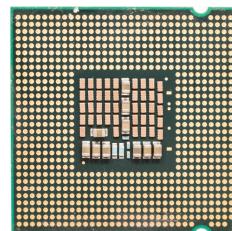
Etiam vel erat nec dui aliquet vulputate sed quis nulla. Donec eget ultricies magna, eu dignissim elit. Nullam sed uma nec nisl rhoncus ullamcorper placerat et enim. Integer varius ornare libero quis consequat. Etiam a purus suscipit, accumsan nibh vel, posuere ipsum. Nulla nec tempor nibh, id venenatis lectus. Duis lobortis id urna eget tincidunt.

Praesent nec libero metus. Praesent at turpis placerat, congue ipsum eget, scelerisque justo. Ut volutpat, massa ac lacinia cursus, nisl dui volutpat arcu, quis interdum sapien turpis in tellus. Suspendisse potenti. Vestibulum pharetra justo massa, ac venenatis mi condimentum nec. Proin viverra tortor non orci suscipit rutrum. Phasellus sit amet euismod diam. Nullam convallis nunc sit amet diam suscipit dapibus. Integer porta hendrerit nunc. Quisque pharetra congue porta. Suspendisse vestibulum sed mi in euismod. Etiam a purus suscipit, accumsan nibh vel, posuere ipsum. Nulla nec tempor nibh, id venenatis lectus. Duis lobortis id urna eget tincidunt.



Overall architecture

SQL



Language

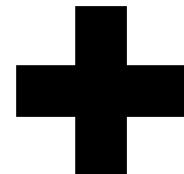
Model

Compute

Storage

Data model

What data **looks like**



What you can **do** with it

Analogy: a very simple "model of cooking"

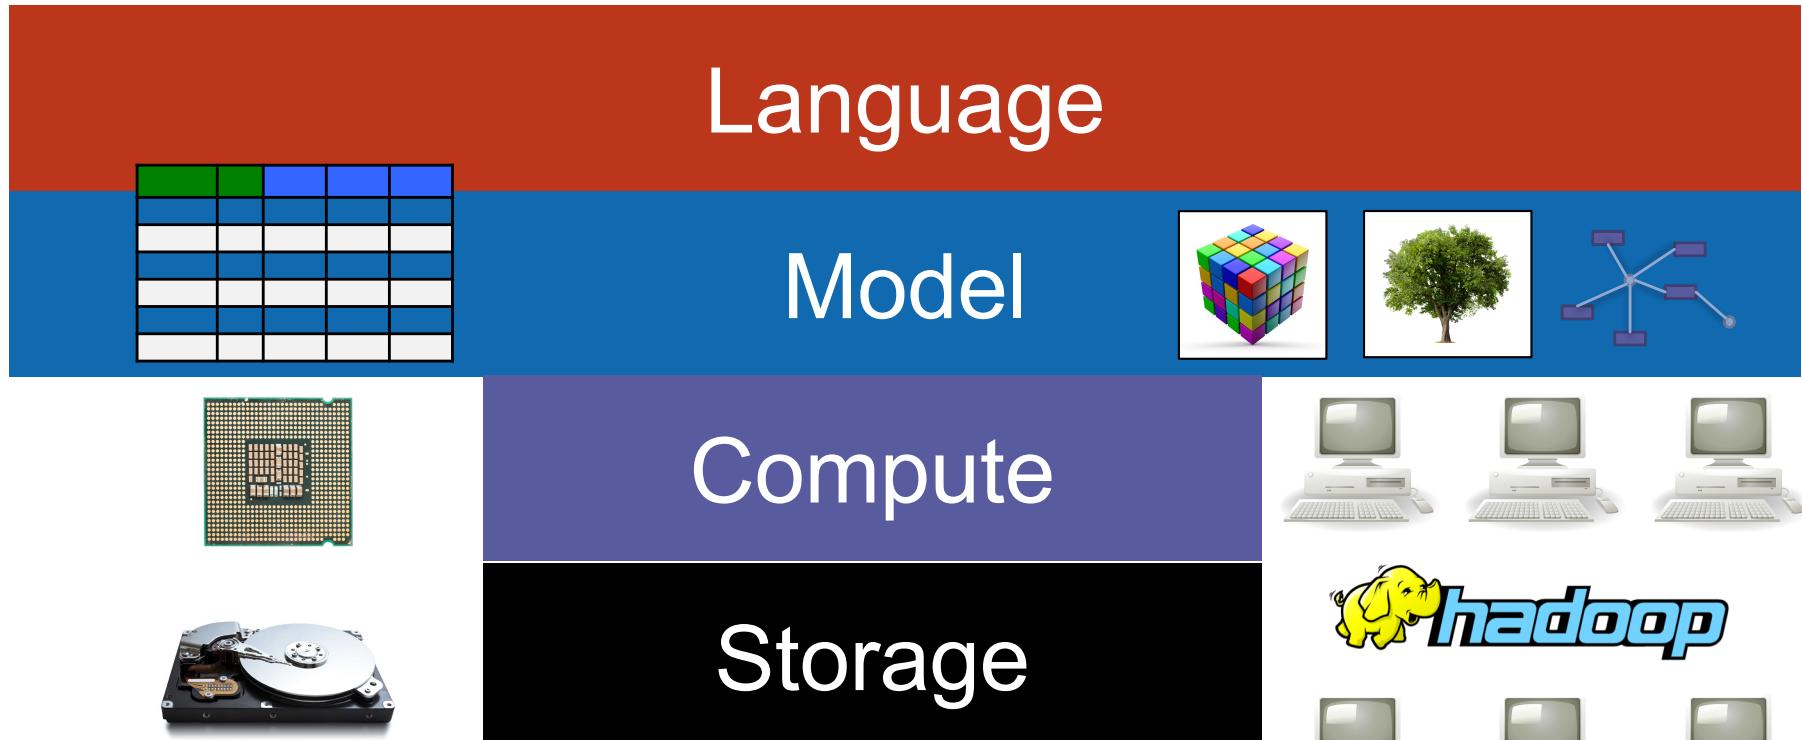
What food **looks like**

→ Something that can be eaten

What you can **do** with it

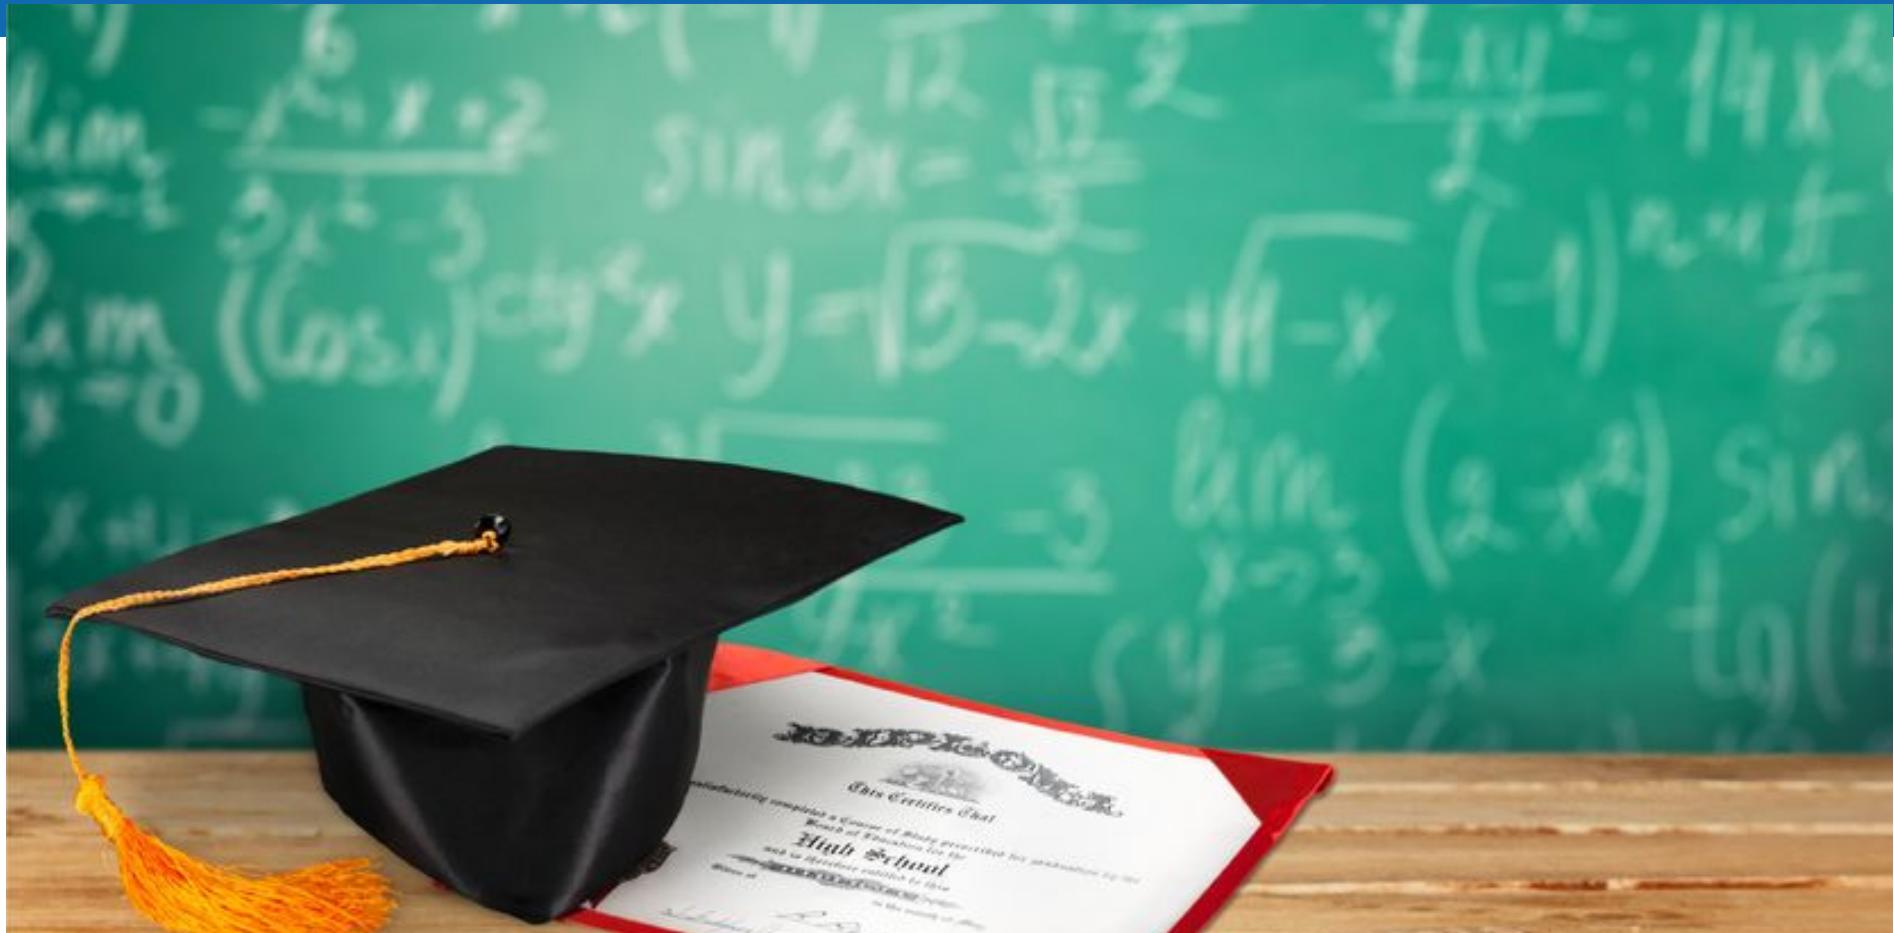
→ You can **mix, pour, break, dry, bake, grill, melt, mince, mix...**

Overall architecture



Old

New



Take-away Concepts

Take-away Concepts

Table
Collection

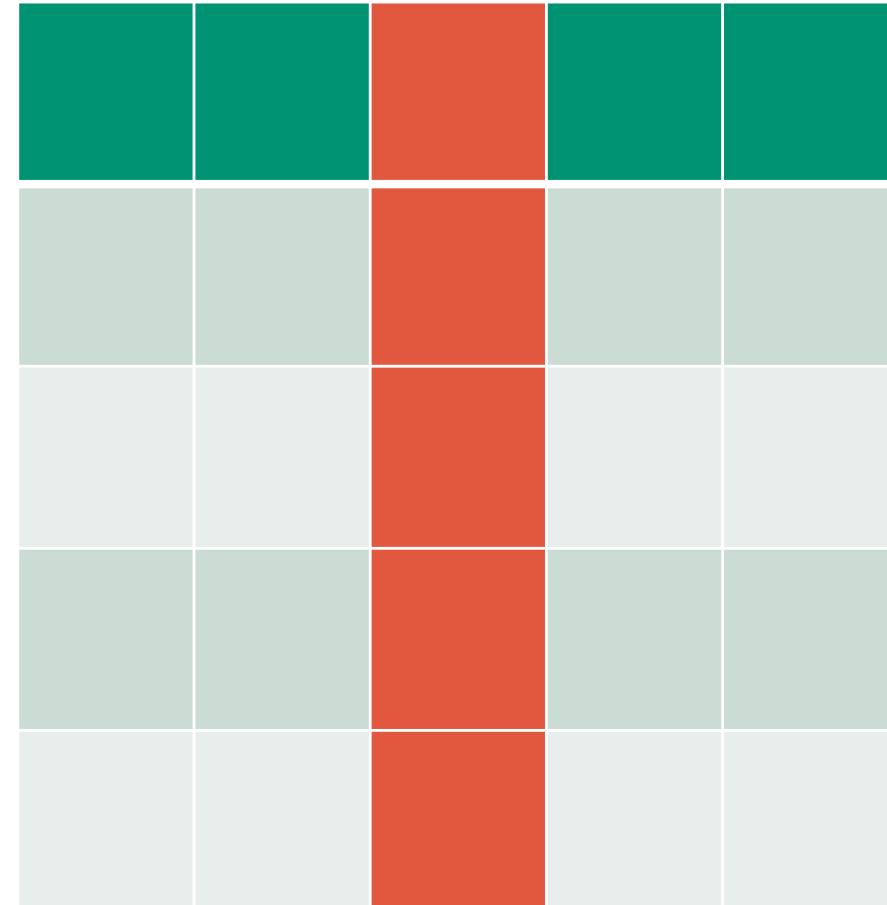
Take-away Concepts

Attribute

Column

Field

Property



Take-away Concepts

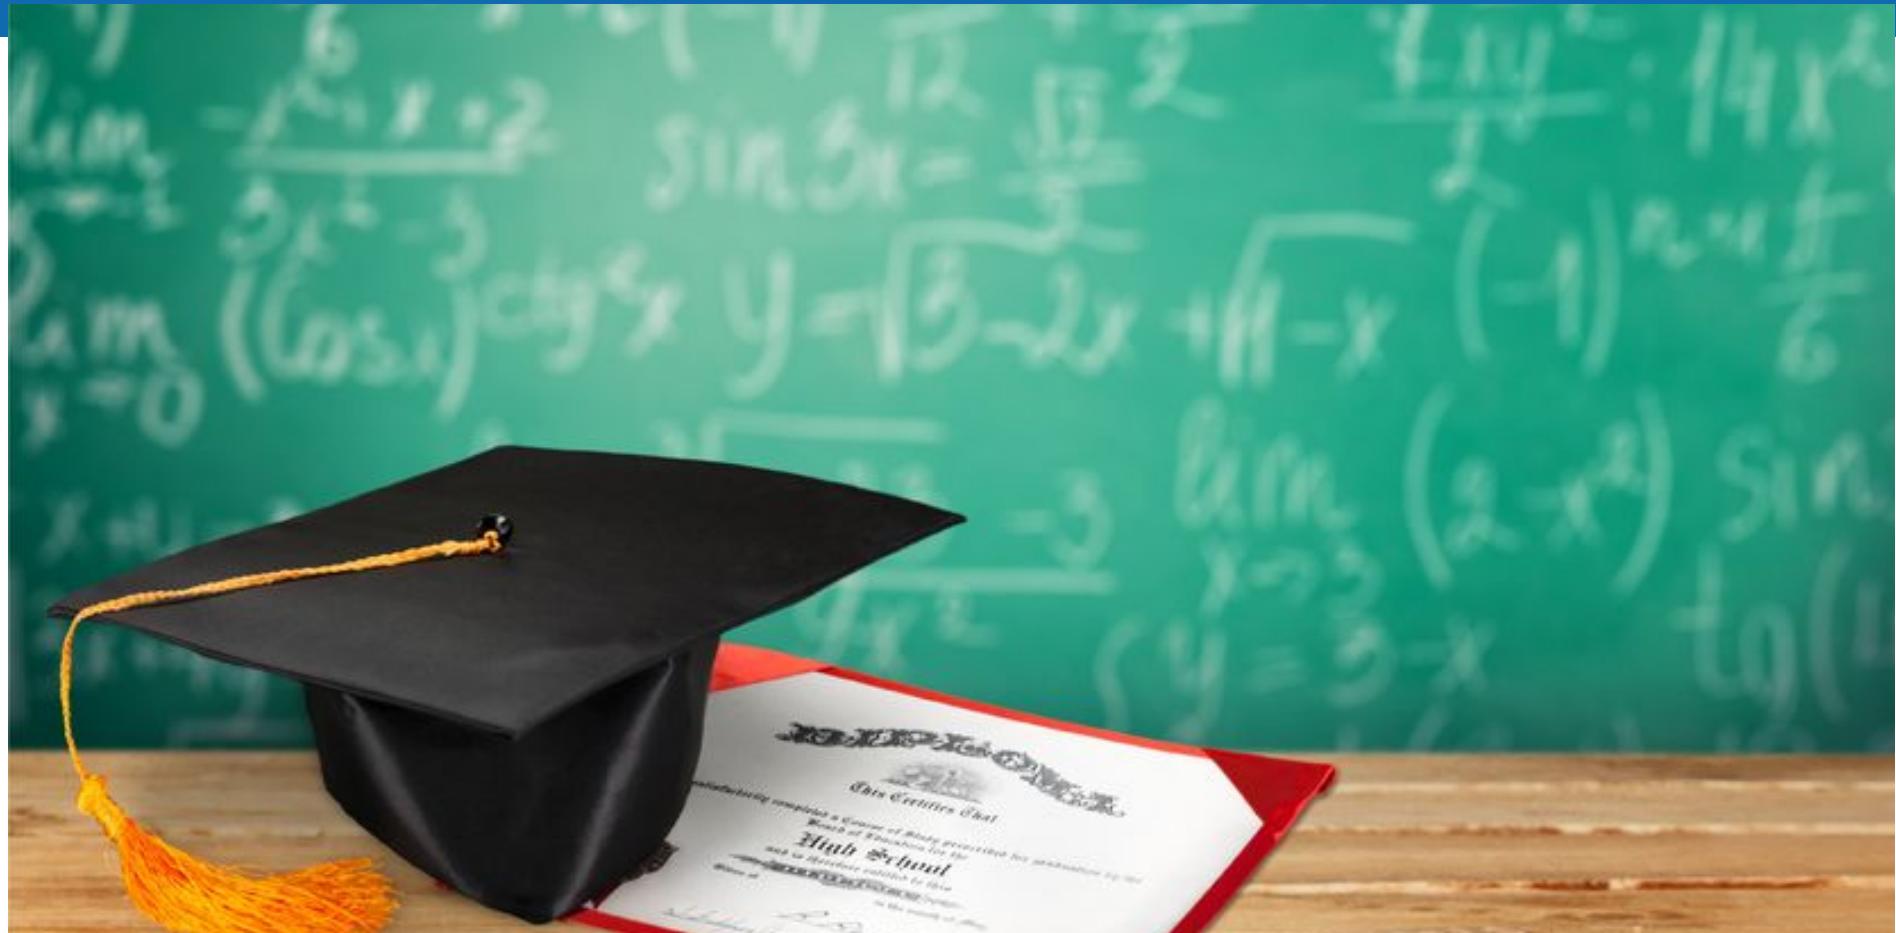
Primary Key

Row ID

Name

Take-away Concepts

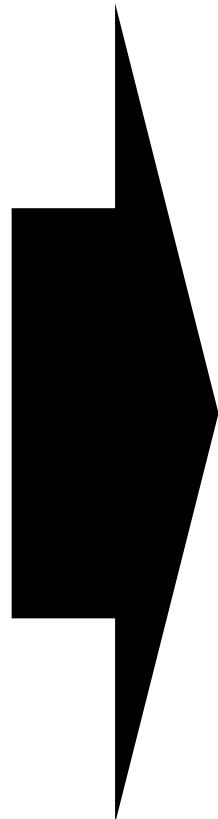
Row
Business Object
Item
Entity
Document
Record



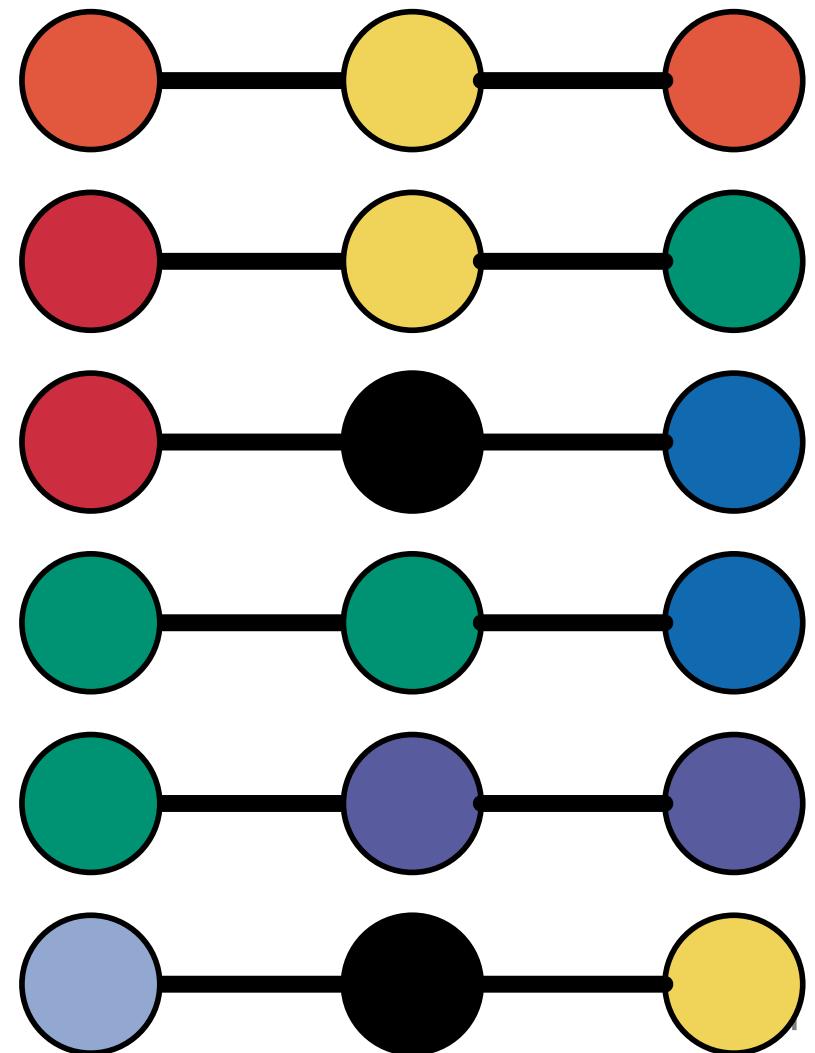
Relational Algebra

Table as a relation

1	2	3
Orange	Yellow	Orange
Red	Yellow	Green
Red	Black	Blue
Green	Green	Blue
Green	Purple	Purple
Blue	Black	Yellow



$$R \subseteq \mathcal{D}_1 \times \mathcal{D}_2 \times \mathcal{D}_3$$



Relations (the math)

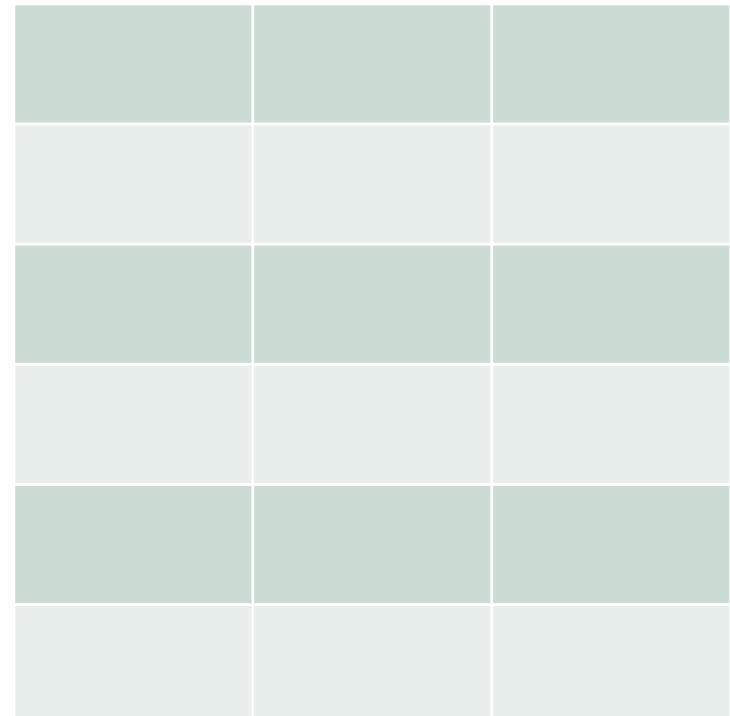
A relation R is made of

1. A **set of attributes**
2. An **extension** (set of tuples)

$$Attributes_R \subseteq \mathbb{S}$$



$$Extension_R \subseteq \mathbb{S} \nrightarrow \mathbb{V}$$



Tuple: example

Name \mapsto Einstein

First name \mapsto Albert

Physicist \mapsto true

Year \mapsto 1905

S

V

Tuple: more intuitive display

S	Name	First name	Physicist	Year
V	Einstein	Albert	true	1905

(The order is irrelevant)

Rule #1: Tabular integrity

Name	First name	Physicist	Year
------	------------	-----------	------



Name	First name	Physicist	Year
Einstein	Albert	true	1905



Name	First name	Physicist	Year
Turing	Alan	false	1936



Name	First name	Physicist	Year
Gödel	Kurt	false	1931

Rule #1: Tabular integrity

Name	First name	Physicist	Year
Einstein	Albert	true	1905
Turing	Alan	false	1936
Gödel	Kurt	false	1931

Rule #1: Broken tabular integrity

	Name	First name	Physicist	Year
Country	Einstein	Albert	true	1905
A	Gödel	Kurt		1931

Rule #2: Atomic integrity (1st normal form)

Legi	Name	Lecture ID	Lecture Name	City	State	PLZ
32-000-000	Alan Turing	xxx-xxxx-xxX	Cryptography	Bletchley Park	UK	MK3 6EB
32-000-000	Alan Turing	263-3010-00L	Big Data	Bletchley Park	UK	MK3 6EB
62-000-000	Georg Cantor	263-3010-00L	Big Data	Pfäffikon	SZ	8808
62-000-000	Georg Cantor	123-4567-89L	Set theory	Pfäffikon	SZ	8808
25-000-000	Felix Bloch	123-4567-89L	Set theory	Pfäffikon	ZH	8330

Rule #2: Broken atomic integrity

Legi	Name	Lecture		City	State	PLZ
32-000-000	Alan Turing	Lecture ID	Lecture Name	Bletchley Park	UK	MK3 6EB
		xxx-xxxx-xxX	Cryptography			
		263-3010-00L	Big Data			
62-000-000	Georg Cantor	Lecture ID	Lecture Name	Pfäffikon	SZ	8808
		263-3010-00L	Big Data			
		123-4567-89L	Set theory			
25-000-000	Felix Bloch	Lecture ID	Lecture Name	Pfäffikon	ZH	8330
		123-4567-89L	Set theory			

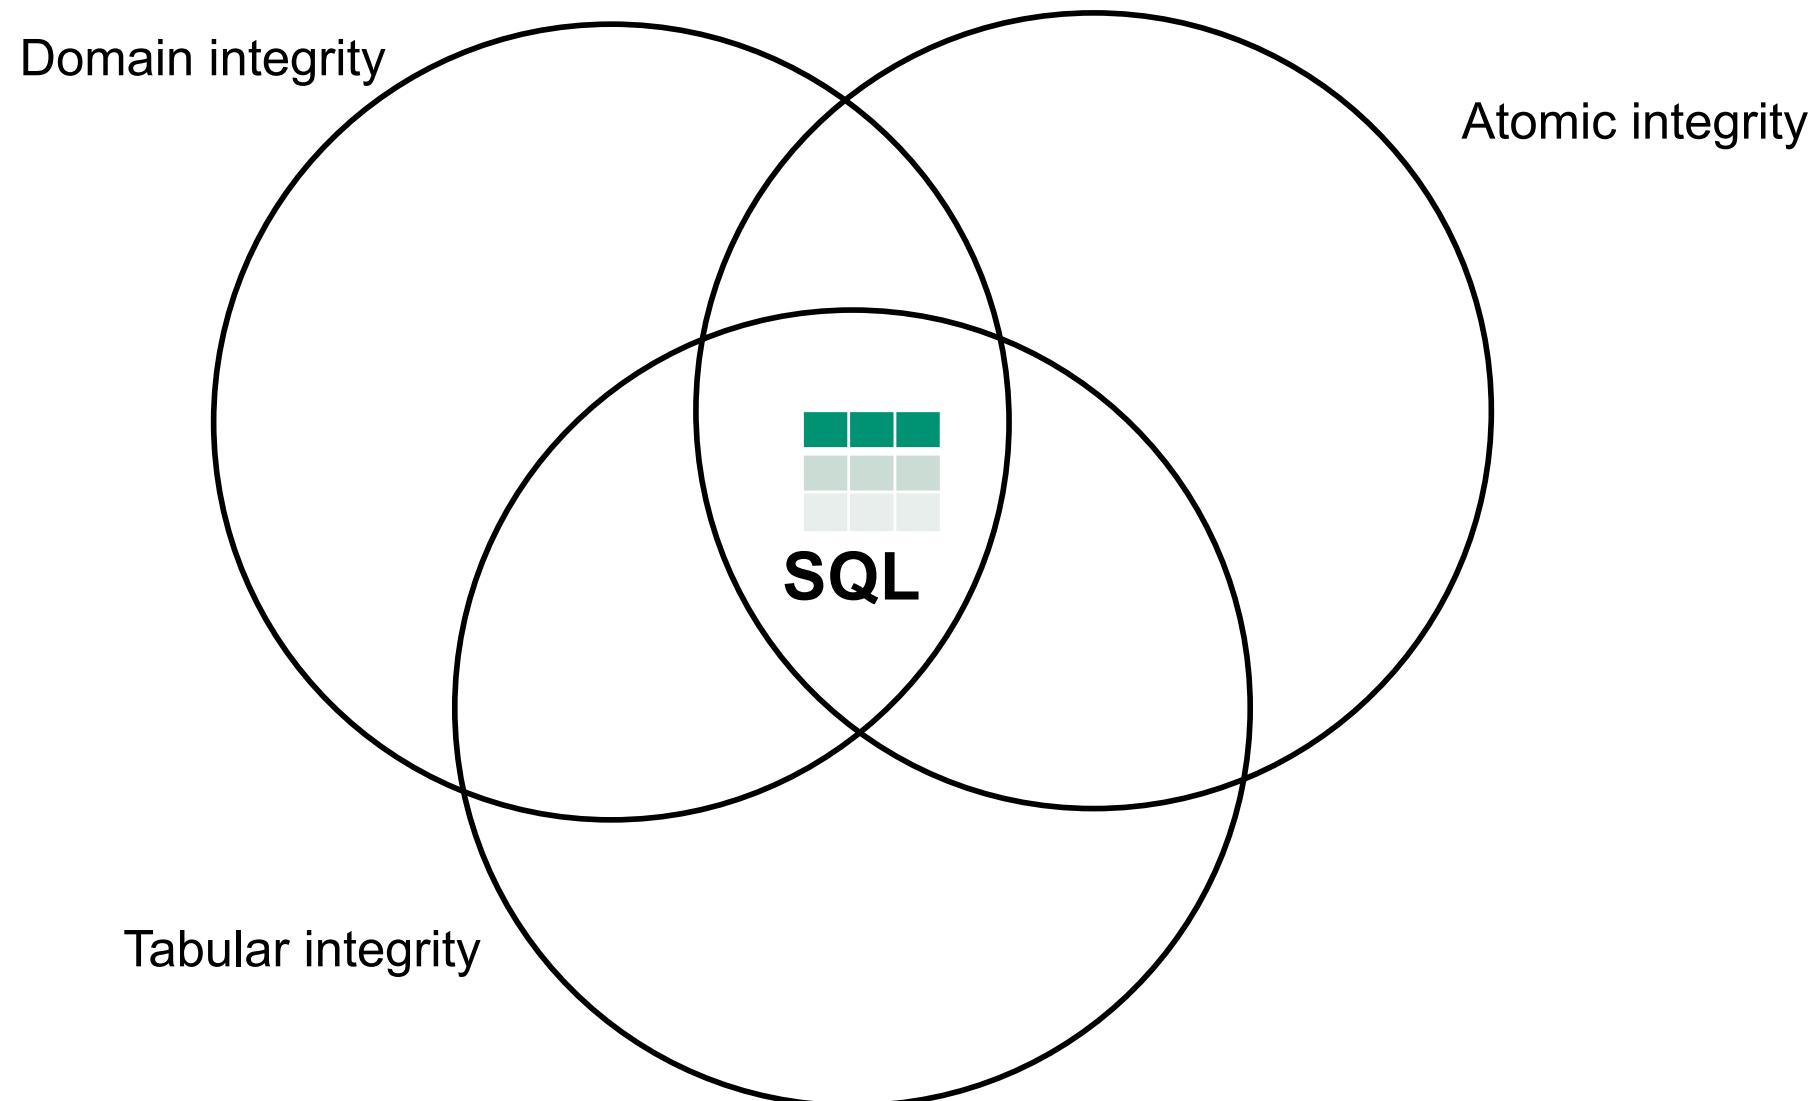
Rule #3: Domain integrity

Name	First name	Physicist	Year
String	String	Boolean	Integer
Einstein	Albert	true	1905
Turing	Alan	false	1936
Gödel	Kurt	false	1931

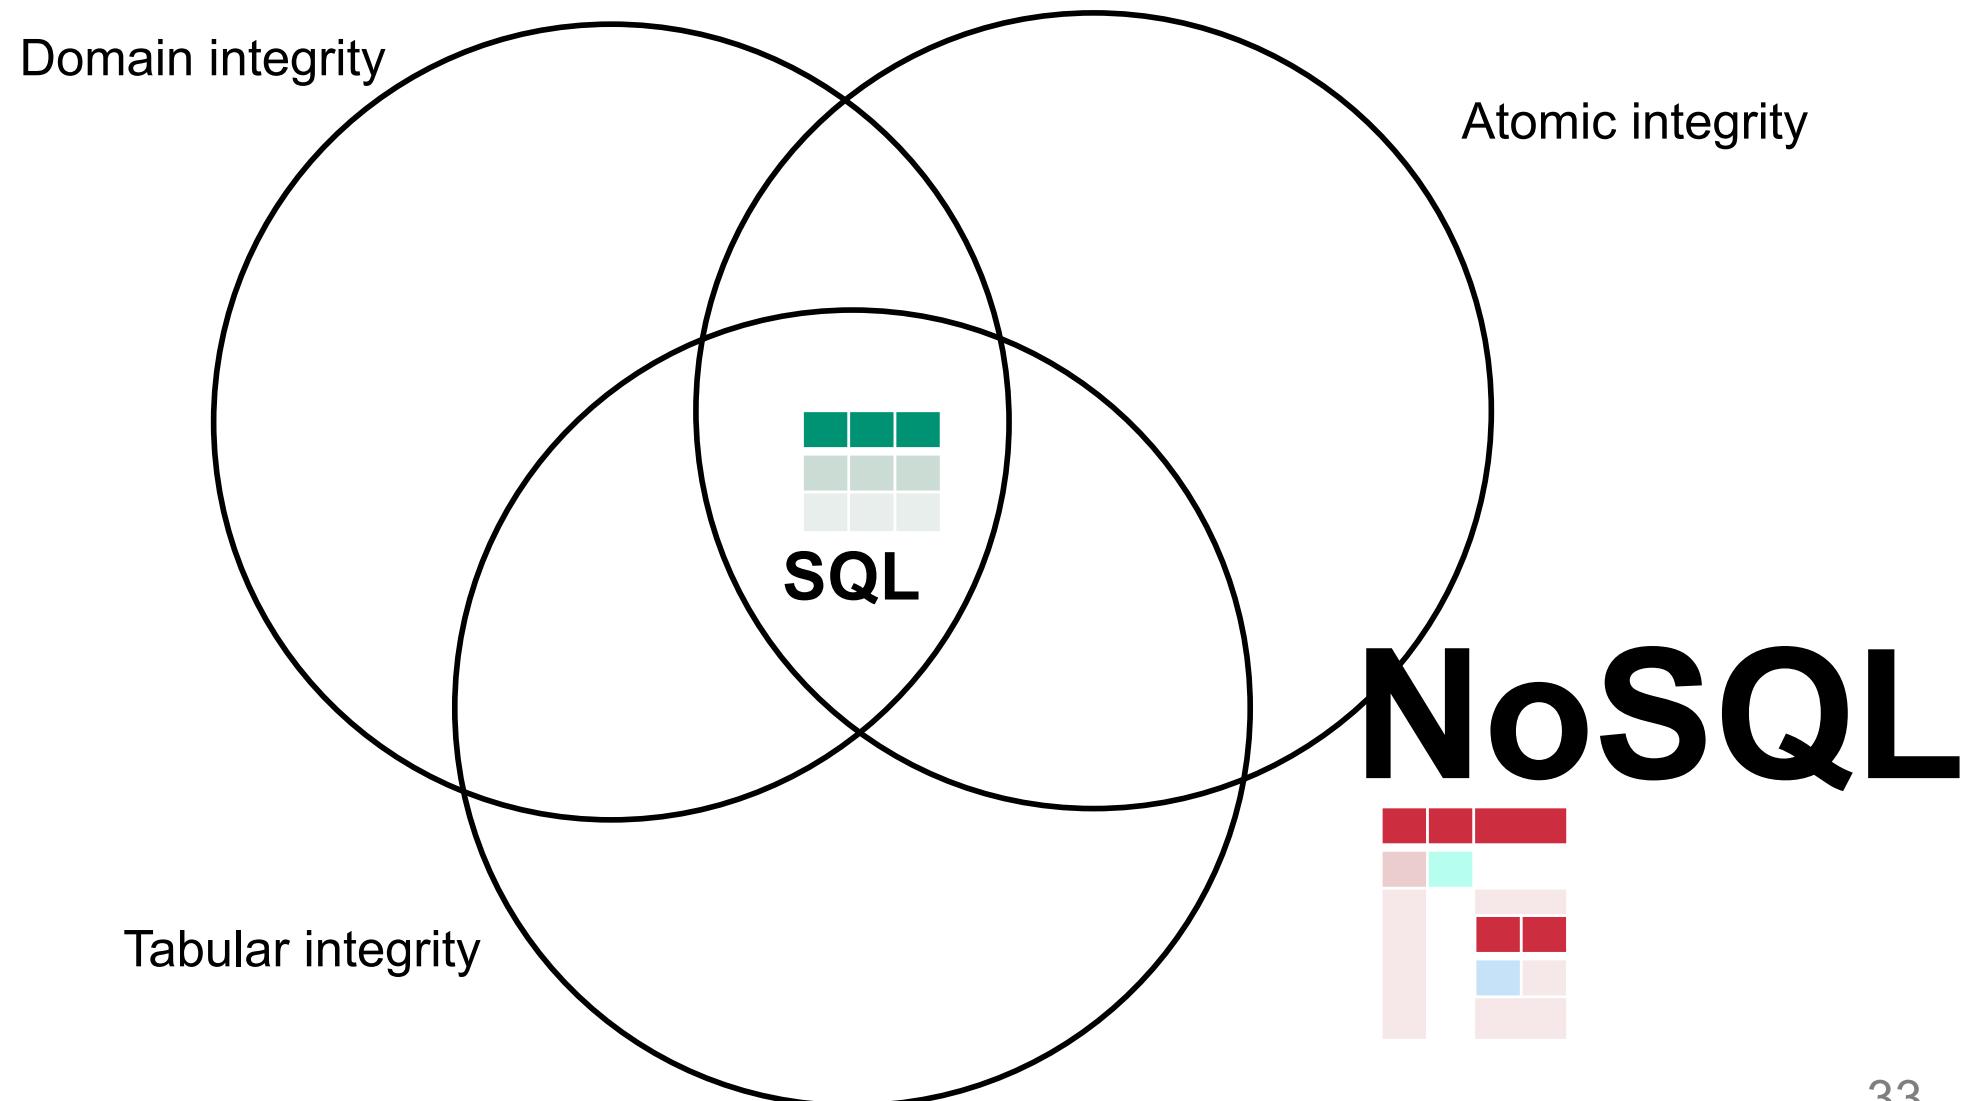
Rule #3: Broken domain integrity

Name	First name	Physicist	Year
String	String	Boolean	Integer
Einstein	Albert	true	1905
Turing	Alan	0	1936
Gödel	Kurt	false	thirty-one

From SQL to NoSQL



From SQL to NoSQL

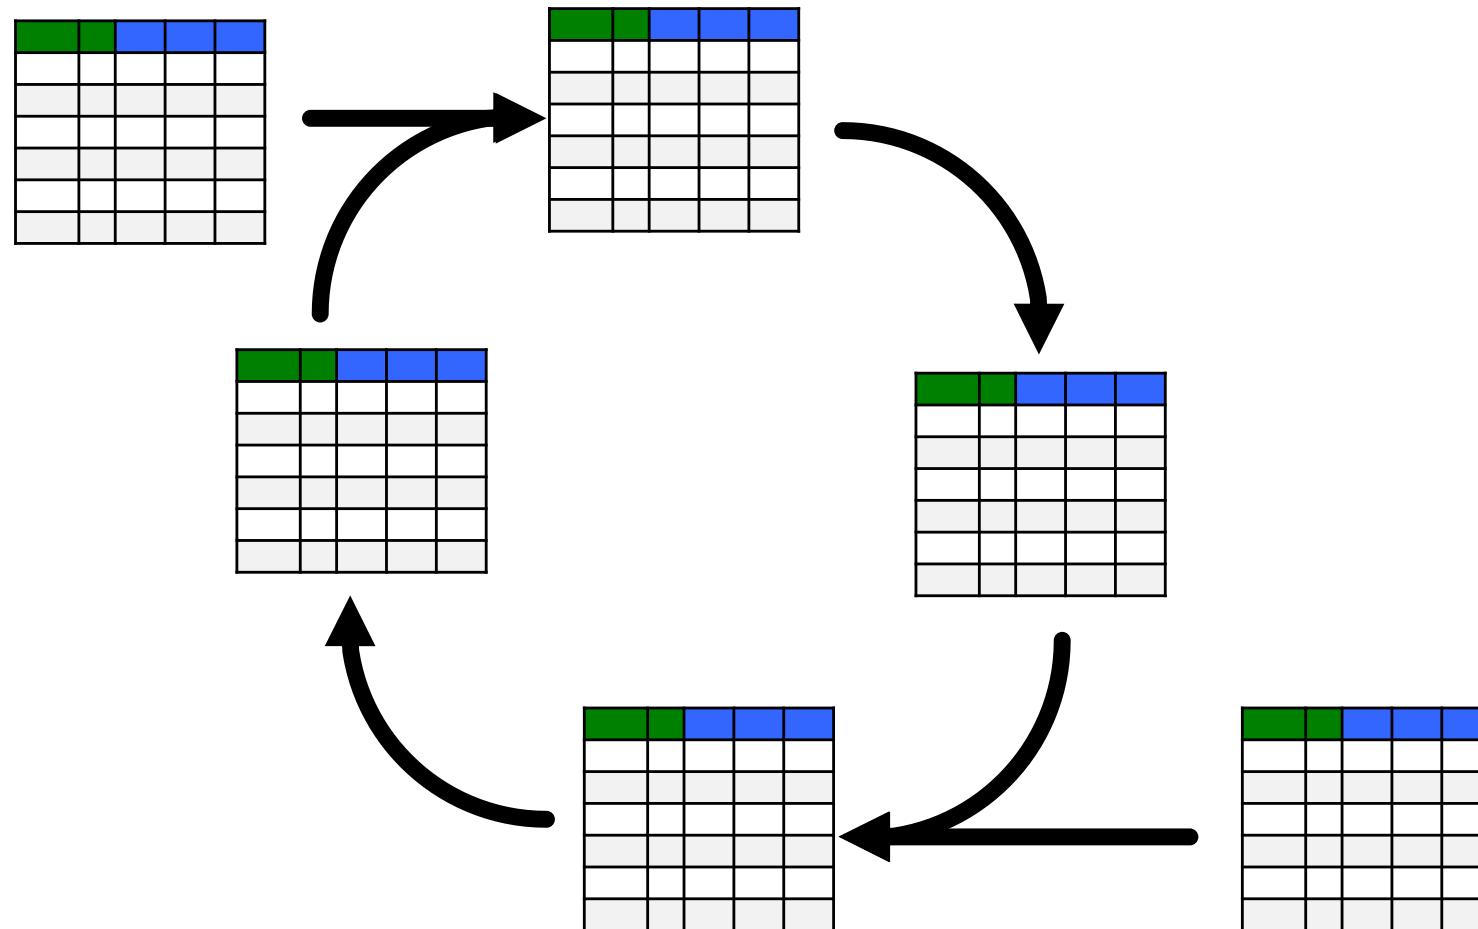


From SQL to NoSQL



SQL

Relational Algebra



Summary of relational queries

Union
Intersection
Subtraction

Set queries

Selection
Projection

Filter queries

Relation renaming
Attribute renaming

Renaming queries

Cartesian product
Natural join
Theta join

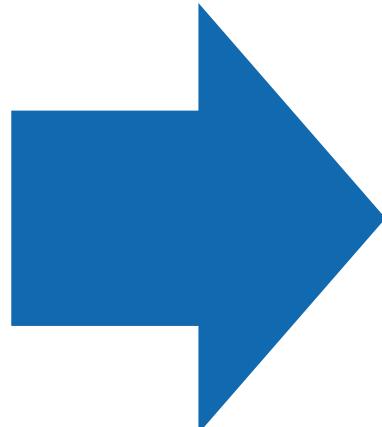
Binary queries

Selection

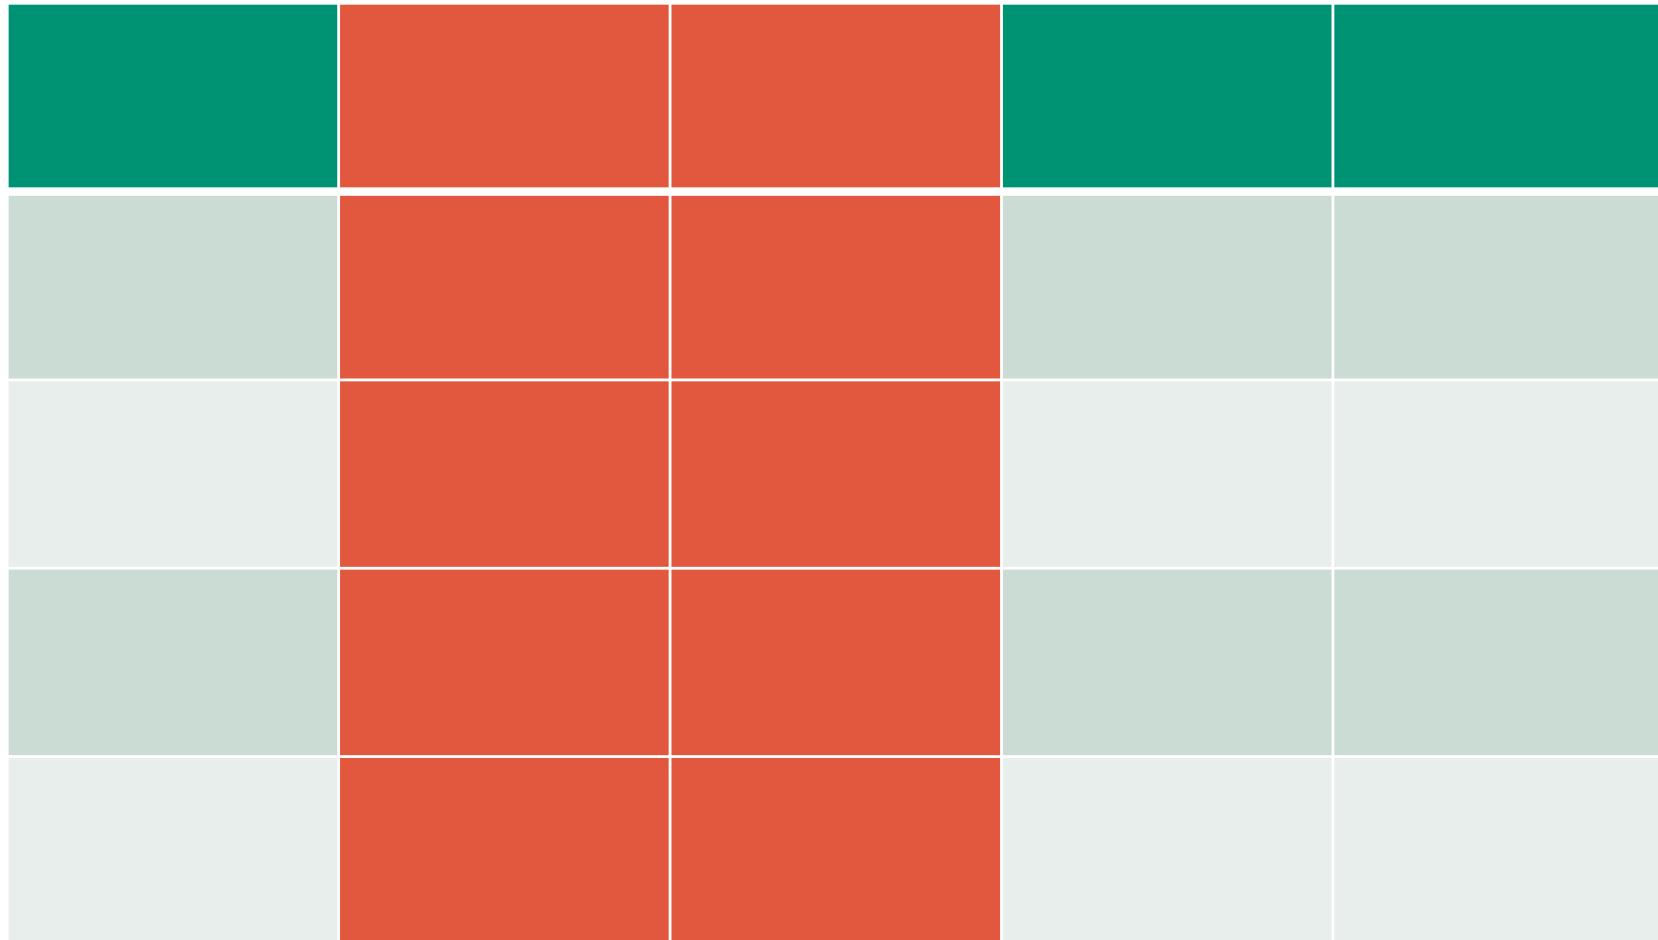
Selection

R		
A	B	C
string	integer	boolean
foo	1	true
bar	2	false
foo	3	false
foobar	4	true



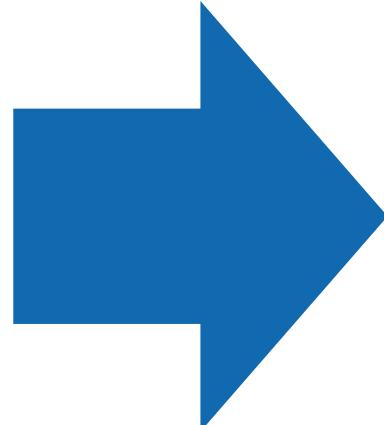
S		
A	B	C
string	integer	boolean
foo	1	true
bar	2	false

Projection



Projection

R		
A	B	C
string	integer	boolean
foo	1	true
bar	2	false
foo	3	false
foobar	4	true



S	
A	C
string	boolean
foo	true
bar	false
foo	false
foobar	true

Grouping

A				
B				
C				
D				

Grouping

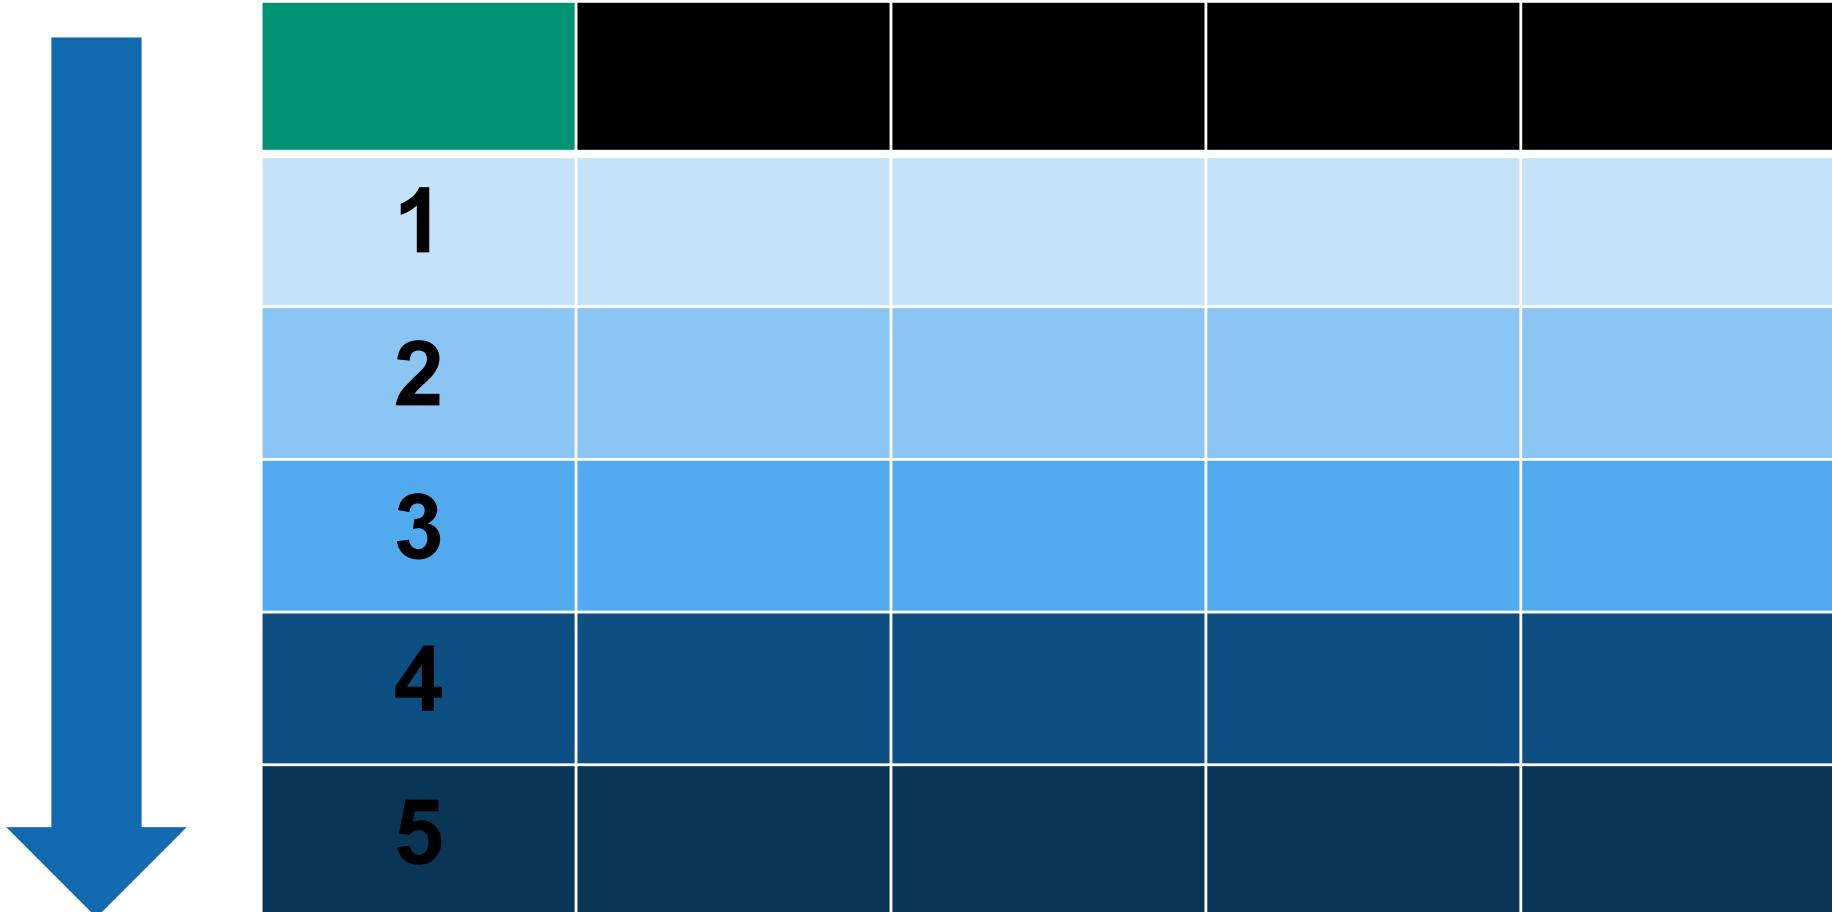
R	
G	A
string	integer
foo	19
bar	28
bar	265
foo	4
foobar	54
foo	46
bar	245
foobar	3456
bar	139

R	
G	A
string	integer
foo	19
foo	4
foo	46
bar	28
bar	265
bar	245
bar	139
foobar	54
foobar	3456

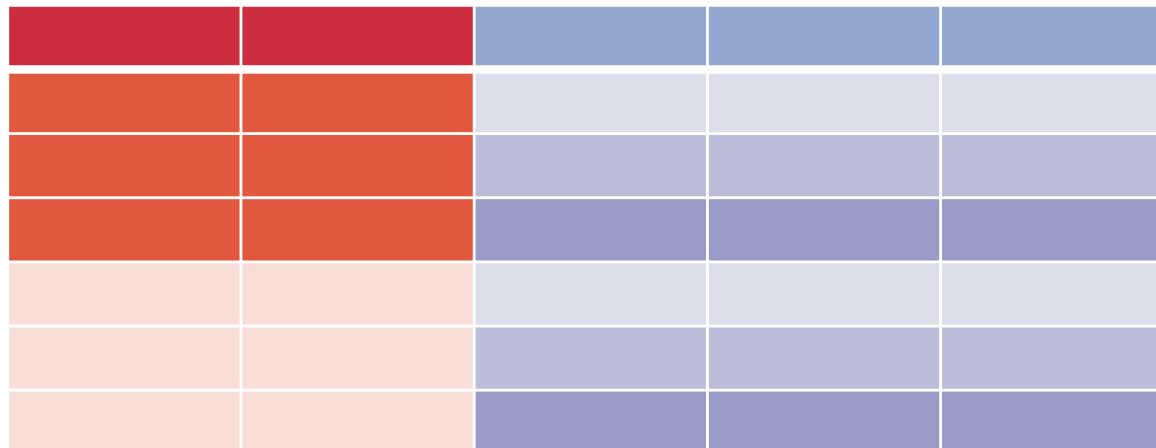
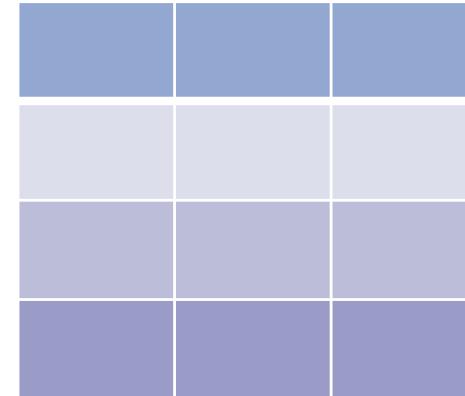
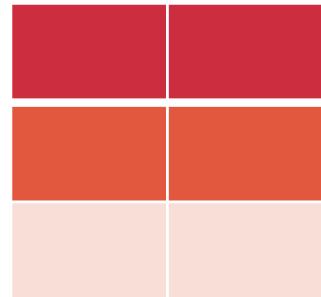
R	
G	A
string	integer
foo	19
bar	4
bar	46
bar	28
bar	265
bar	245
bar	139
foobar	54
foobar	3456

R	
G	A
string	integer
foo	69
bar	677
foobar	3510

Sorting



Cartesian product



Cartesian product

R			T				
A	B	C	A	B	C	D	E
string	integer	boolean	string	integer	boolean	string	integer
foo	1	true	foo	1	true	foo	1
bar	2	false	foo	1	true	bar	2
S			foo	1	true	foo	3
D	E	bar	2	false	foo	1	
string	integer	bar	2	false	bar	2	
foo	1	bar	2	false	foo	3	
bar	2	foo	3				

Join

	A	
B		



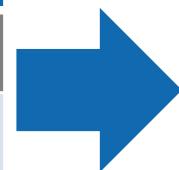
A		
B		
A		



	A		
	A		
B			

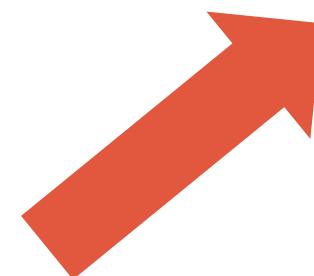
Join

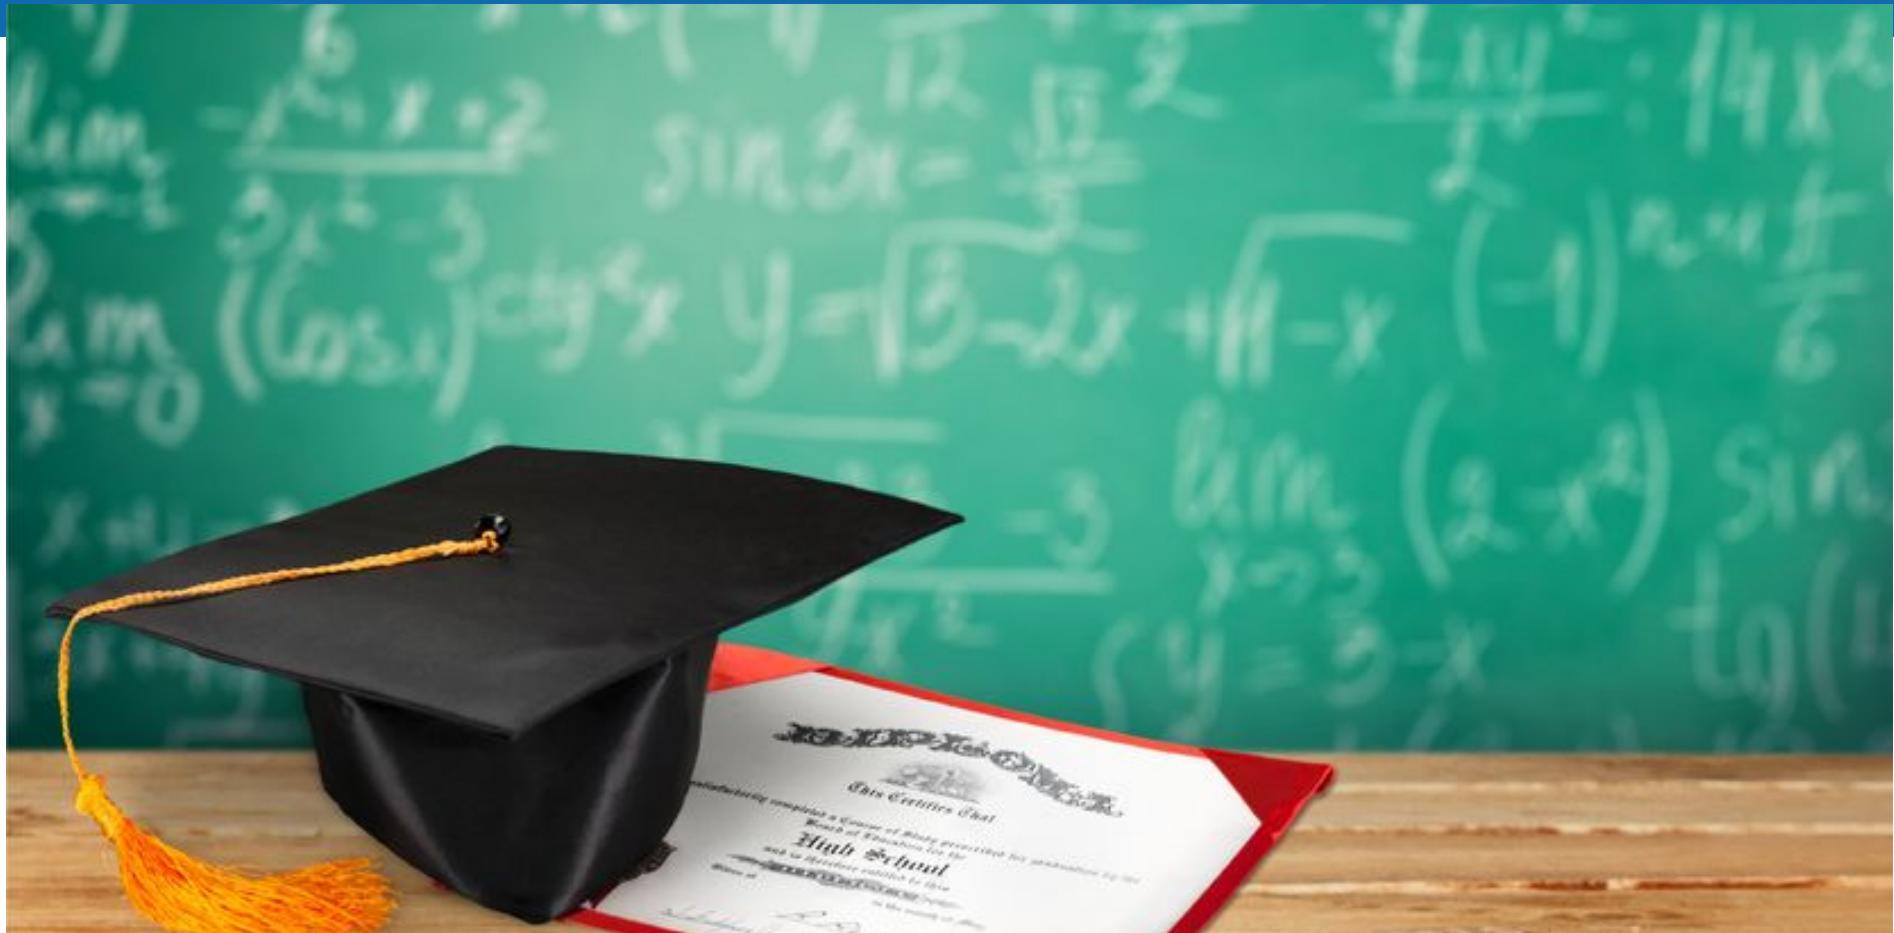
R		
A	B	C
string	integer	boolean
foo	1	true
bar	2	false



T			
A	B	C	D
string	integer	boolean	string
foo	1	true	bar
bar	2	false	bar

S	
D	B
string	integer
bar	1
bar	2
foo	3





SQL Brush-Up

SQL History



Don Chamberlin



Raymond Boyce

The early days (early 1970s)



Almaden (San Jose)

First commercial relational database

System R

+

SEQUEL

First commercial relational query language



Pratt & Whitney

A United Technologies Company

First customer (1977)

Public availability

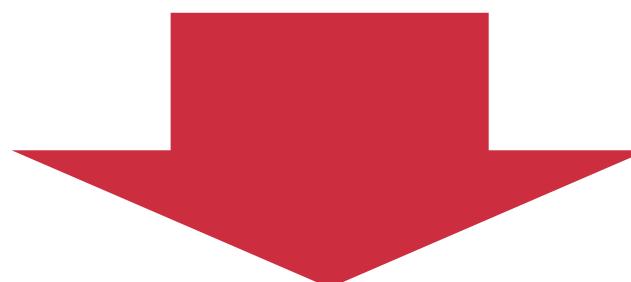
1977

Software Development Laboratories



1979

Relational Software



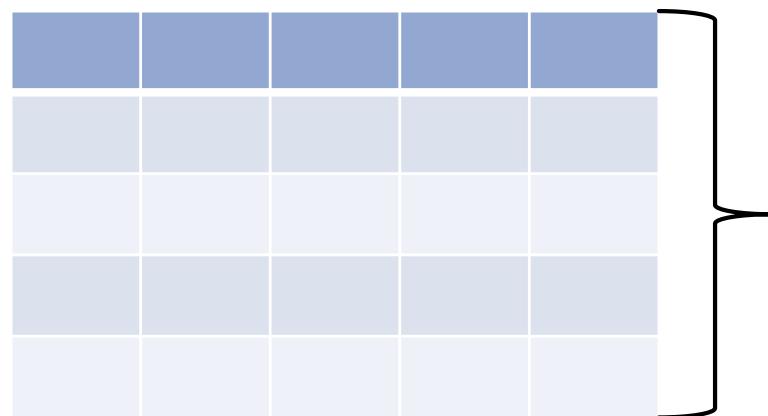
1982

ORACLE

SEQUEL

Structured English QUERy Language

Declarative language

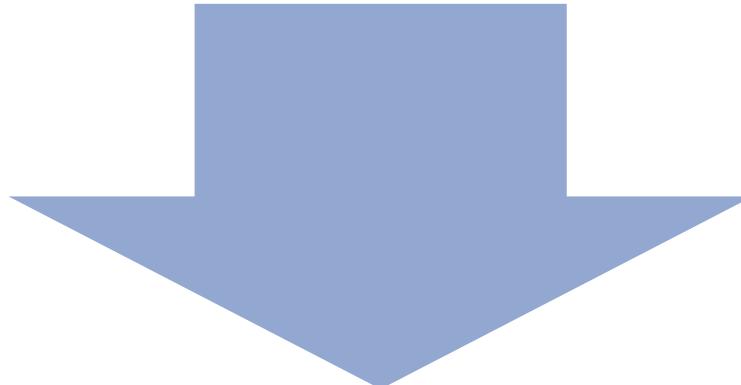


Set-based
(Manipulates entire relations
with a single command)

Renaming



SEQUEL



(Trademark issue)

SQL

ESS-kew-EL

or

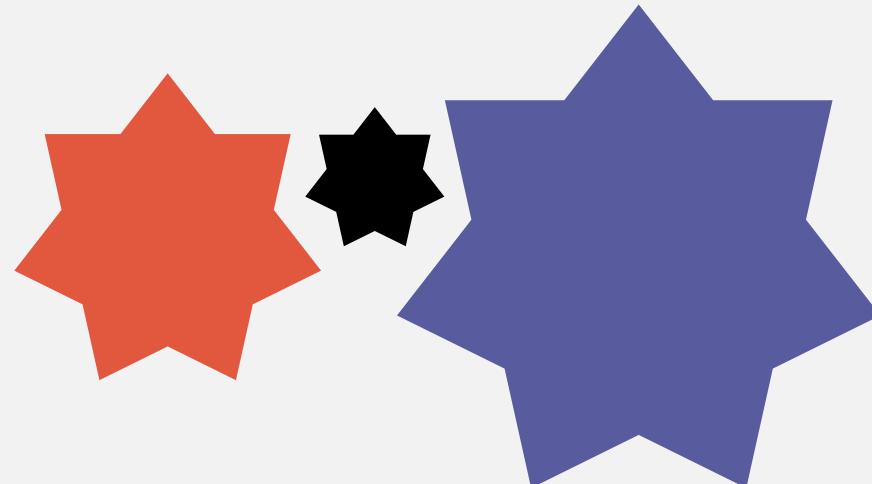
SEE-kwəl

SQL is a declarative language

Logical model

"What, not how"

Physical execution

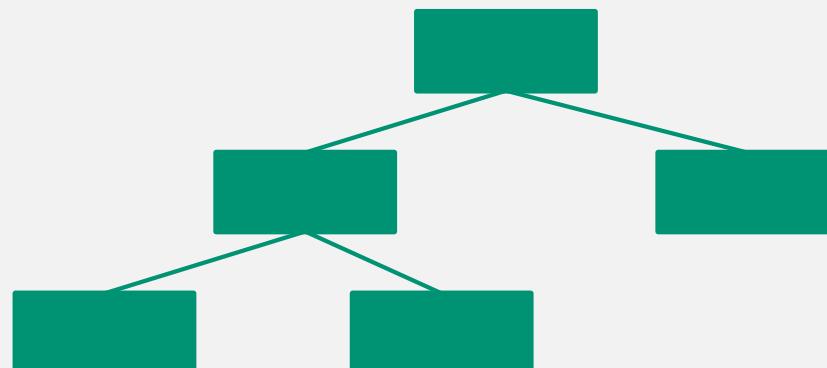


SQL is a declarative language

Logical model

"What, not how"

Query Plan

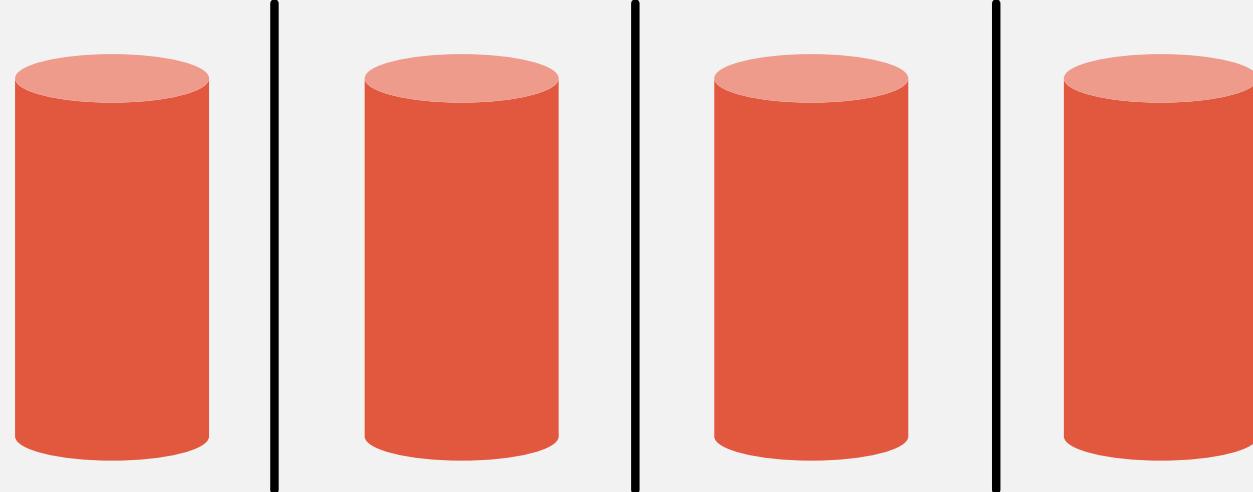


SQL is a declarative language

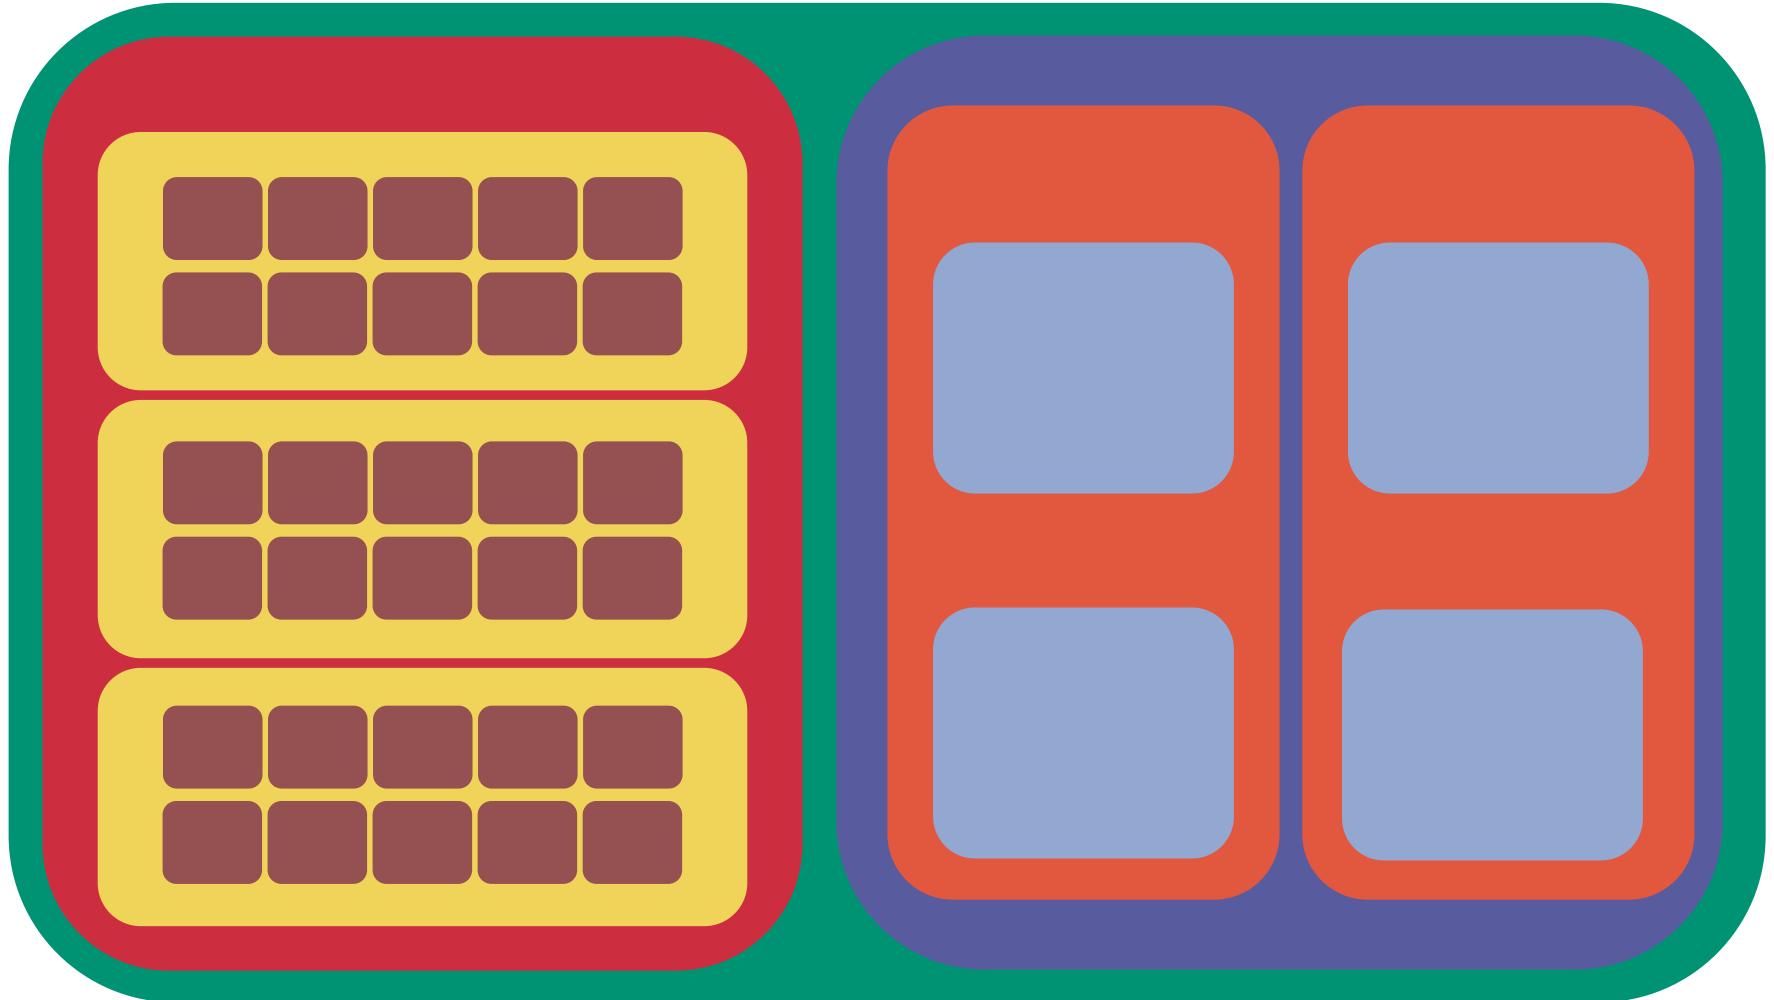
Logical model

"What, not how"

Parallelism



SQL is a functional language



"Expressions that nest"

A selecting query

persons					
name	middle_initial	last_name	birth_date	gender	passport_scan
varchar(30)	char(1)	text	date	boolean	bytea
James	T	Kirk	2233-03-22	TRUE	AD10E7
Beverly	C	Crusher	2324-10-13	FALSE	AD234F7
Spock	NULL	NULL	2230-01-06	TRUE	AD234F7

```
SELECT *
  FROM persons
 WHERE last_name = 'Crusher'
```

persons					
name	middle_initial	last_name	birth_date	gender	passport_scan
varchar(30)	char(1)	text	date	boolean	bytea
Beverly	C	Crusher	2324-10-13	FALSE	AD234F7

A projecting query

persons					
name	middle_initial	last_name	birth_date	gender	passport_scan
varchar(30)	char(1)	text	date	boolean	bytea
James	T	Kirk	2233-03-22	TRUE	AD10E7
Beverly	C	Crusher	2324-10-13	FALSE	AD234F7
Spock	NULL	NULL	2230-01-06	TRUE	AD234F7

SELECT name, birth_date
FROM persons

name	birth_date
varchar(30)	date
James	2233-03-22
Beverly	2324-10-13
Spock	2230-01-06

A renaming query

persons					
name	middle_initial	last_name	birth_date	gender	passport_scan
varchar(30)	char(1)	text	date	boolean	bytea
James	T	Kirk	2233-03-22	TRUE	AD10E7
Beverly	C	Crusher	2324-10-13	FALSE	AD234F7
Spock	NULL	NULL	2230-01-06	TRUE	AD234F7

```
SELECT name AS who,  
       birth_date AS when  
  FROM persons
```

who	when
varchar(30)	date
James	2233-03-22
Beverly	2324-10-13
Spock	2230-01-06

A sorting query

persons					
name	middle_initial	last_name	birth_date	gender	passport_scan
varchar(30)	char(1)	text	date	boolean	bytea
James	T	Kirk	2233-03-22	TRUE	AD10E7
Beverly	C	Crusher	2324-10-13	FALSE	AD234F7
Spock	NULL	NULL	2230-01-06	TRUE	AD234F7

```
SELECT *
  FROM persons
 ORDER BY birth_date
```

persons					
name	middle_initial	last_name	birth_date	gender	passport_scan
varchar(30)	char(1)	text	date	boolean	bytea
Spock	NULL	NULL	2230-01-06	TRUE	AD234F7
James	T	Kirk	2233-03-22	TRUE	AD10E7
Beverly	C	Crusher	2324-10-13	FALSE	AD234F7

A grouping query

persons					Grouping column
name	middle_initial	last_name	century	gender	
varchar(30)	char(1)	text	integer	boolean	
James	T	Kirk	23	TRUE	
Beverly	C	Crusher	24	FALSE	
Spock	NULL	NULL	23	TRUE	
Kathryn	NULL	Janeway	24	FALSE	

```
SELECT century, COUNT(*) AS count
FROM persons
GROUP BY century
```

century	count
integer	bigint
23	2
24	2

Post-aggregation selection

persons				
name	middle_initial	last_name	century	captain
varchar(30)	char(1)	text	integer	boolean
James	T	Kirk	23	TRUE
Beverly	C	Crusher	24	FALSE
Jean-Luc	NULL	Picard	24	TRUE
Kathryn	NULL	Janeway	24	TRUE

```
SELECT century AS c
FROM persons
GROUP BY century
HAVING COUNT(*) > 2
```

c
integer
24

Natural join

persons				
name	middle_initial	last_name	century	captain
varchar(30)	char(1)	text	integer	boolean
James	T	Kirk	23	TRUE
Beverly	C	Crusher	24	FALSE
Kathryn	NULL	Janeway	24	TRUE

spaceships			
warp	spaceship_name	last_name	code
numeric	varchar(30)	text	integer
5	USS Enterprise A	Kirk	NCC-1701-A
4	USS Enterprise	Kirk	NCC-1701
9.2	USS Enterprise D	Picard	NCC-1701-D
9.975	USS Voyager	Janeway	NCC-74656

```
SELECT *
FROM persons NATURAL FULL OUTER JOIN spaceships
```

name	middle_initial	last_name	captain	century	warp	spaceship_name	code
varchar(30)	char(1)	text	boolean	integer	numeric	varchar(30)	integer
James	T	Kirk	TRUE	23	5	USS Enterprise A	NCC-1701-A
James	T	Kirk	TRUE	23	4	USS Enterprise	NCC-1701
NULL	NULL	Picard	NULL	NULL	9.2	USS Enterprise D	NCC-1701-D
Beverly	C	Crusher	24	FALSE	NULL	NULL	NULL
Kathryn	NULL	Janeway	TRUE	24	9.975	USS Voyager	NCC-74656

A union query

spaceships1				
warp	name	last_name	code	successor
numeric	varchar(30)	text	varchar	varchar
5	USS Enterprise A	Kirk	NCC-1701-A	NCC-1701-B
6	USS Enterprise B	Kirk	NCC-1701-B	NCC-1701-C
7	USS Enterprise C	Kirk	NCC-1701-C	NCC-1701-D

spaceships2				
warp	name	last_name	code	successor
numeric	varchar(30)	text	varchar	varchar
7	USS Enterprise C	Kirk	NCC-1701-C	NCC-1701-D
4	USS Enterprise	Kirk	NCC-1701	NCC-1701-A
9.2	USS Enterprise D	Picard	NCC-1701-D	NCC-1701-E

`SELECT * FROM spaceships1 UNION spaceships2`

duplicate elimination

warp	name	last_name	code	successor
numeric	varchar(30)	text	varchar	varchar
5	USS Enterprise A	Kirk	NCC-1701-A	NCC-1701-B
6	USS Enterprise B	Kirk	NCC-1701-B	NCC-1701-C
7	USS Enterprise C	Kirk	NCC-1701-C	NCC-1701-D
4	USS Enterprise	Kirk	NCC-1701	NCC-1701-A
9.2	USS Enterprise D	Picard	NCC-1701-D	NCC-1701-E

An intersection query

spaceships1				
warp	name	last_name	code	successor
numeric	varchar(30)	text	varchar	varchar
5	USS Enterprise A	Kirk	NCC-1701-A	NCC-1701-B
6	USS Enterprise B	Kirk	NCC-1701-B	NCC-1701-C
7	USS Enterprise C	Kirk	NCC-1701-C	NCC-1701-D

spaceships2				
warp	name	last_name	code	successor
numeric	varchar(30)	text	varchar	varchar
7	USS Enterprise C	Kirk	NCC-1701-C	NCC-1701-D
4	USS Enterprise	Kirk	NCC-1701	NCC-1701-A

`SELECT * FROM spaceships1 INTERSECT spaceships2`

duplicate elimination

warp	name	last_name	code	successor
numeric	varchar(30)	text	varchar	varchar
7	USS Enterprise C	Kirk	NCC-1701-C	NCC-1701-D

An set subtraction query

spaceships1				
warp	name	last_name	code	successor
numeric	varchar(30)	text	varchar	varchar
5	USS Enterprise A	Kirk	NCC-1701-A	NCC-1701-B
6	USS Enterprise B	Kirk	NCC-1701-B	NCC-1701-C
7	USS Enterprise C	Kirk	NCC-1701-C	NCC-1701-D

spaceships2				
warp	name	last_name	code	successor
numeric	varchar(30)	text	varchar	varchar
7	USS Enterprise C	Kirk	NCC-1701-C	NCC-1701-D
4	USS Enterprise	Kirk	NCC-1701	NCC-1701-A

`SELECT * FROM spaceships1 EXCEPT spaceships2`

duplicate elimination

warp	name	last_name	code	successor
numeric	varchar(30)	text	varchar	varchar
5	USS Enterprise A	Kirk	NCC-1701-A	NCC-1701-B
6	USS Enterprise B	Kirk	NCC-1701-B	NCC-1701-C

Nesting

```
SELECT pname, warp
FROM (
    SELECT persons.name AS pname,
           last_name,
           century,
           spaceships.name AS sname,
           warp
      FROM persons FULL OUTER JOIN spaceships USING last_name
)
```

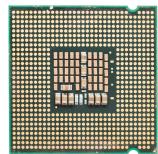
Some terminology

DML: Data Manipulation Language
(Query, insert, remove rows)

Data Schema

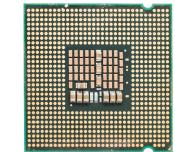
DDL: Data Definition Language
(Create or table/schema, drop it)

Language landscape

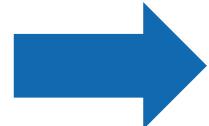


Proto-imperative
language

Language landscape

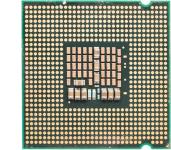


Proto-imperative
language



Imperative
language

Language landscape



Proto-imperative
language

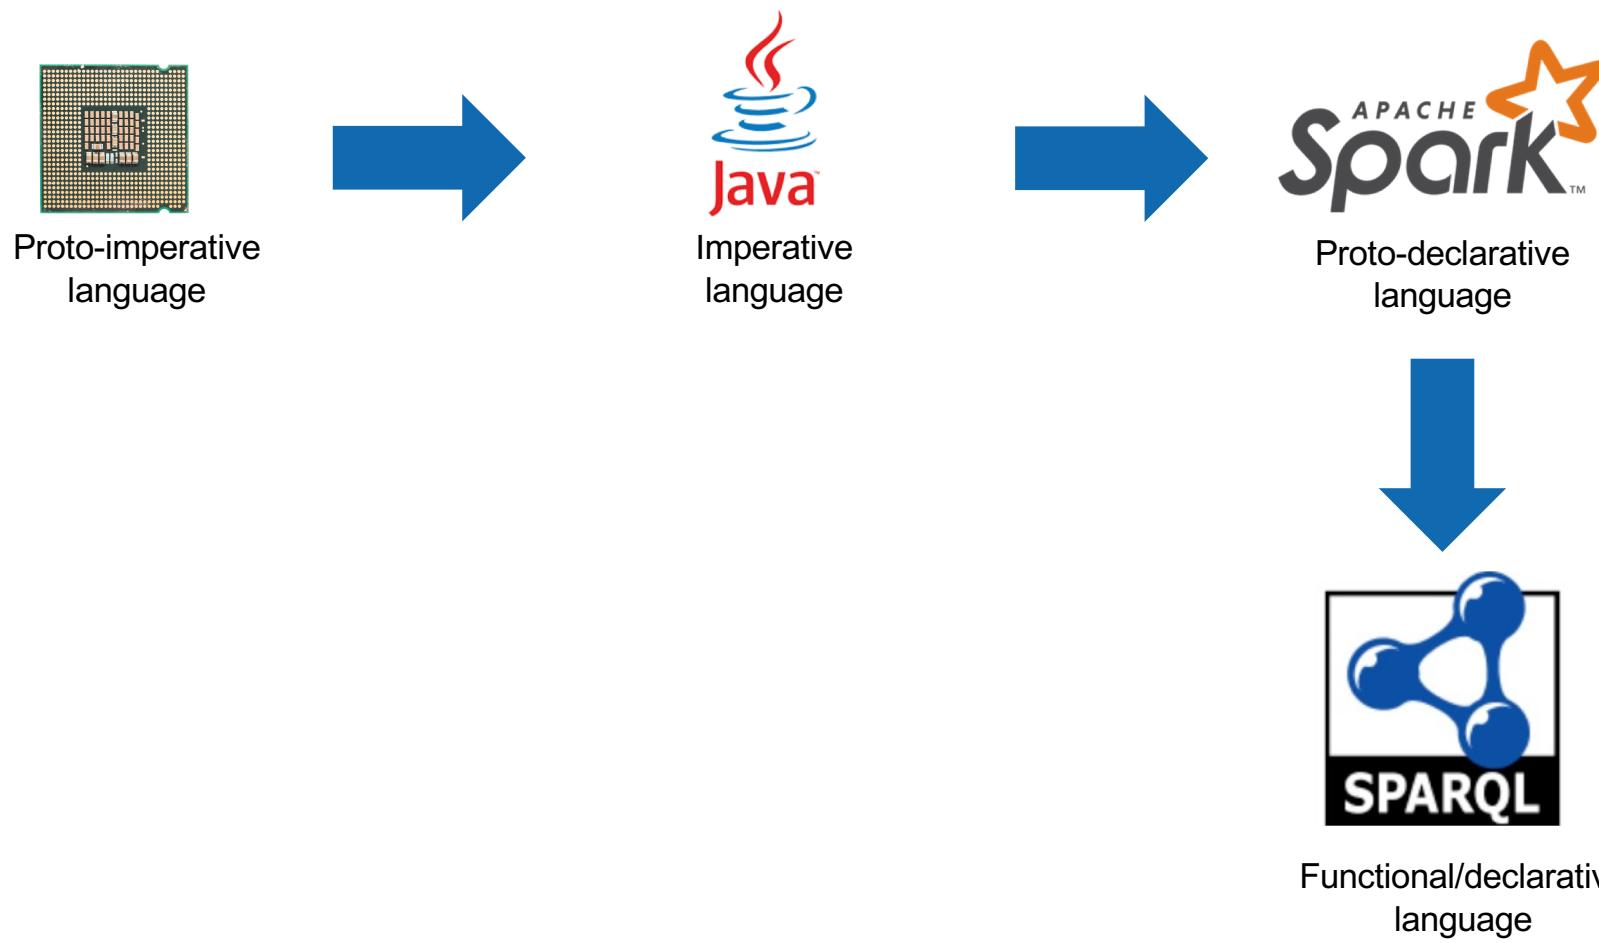


Imperative
language

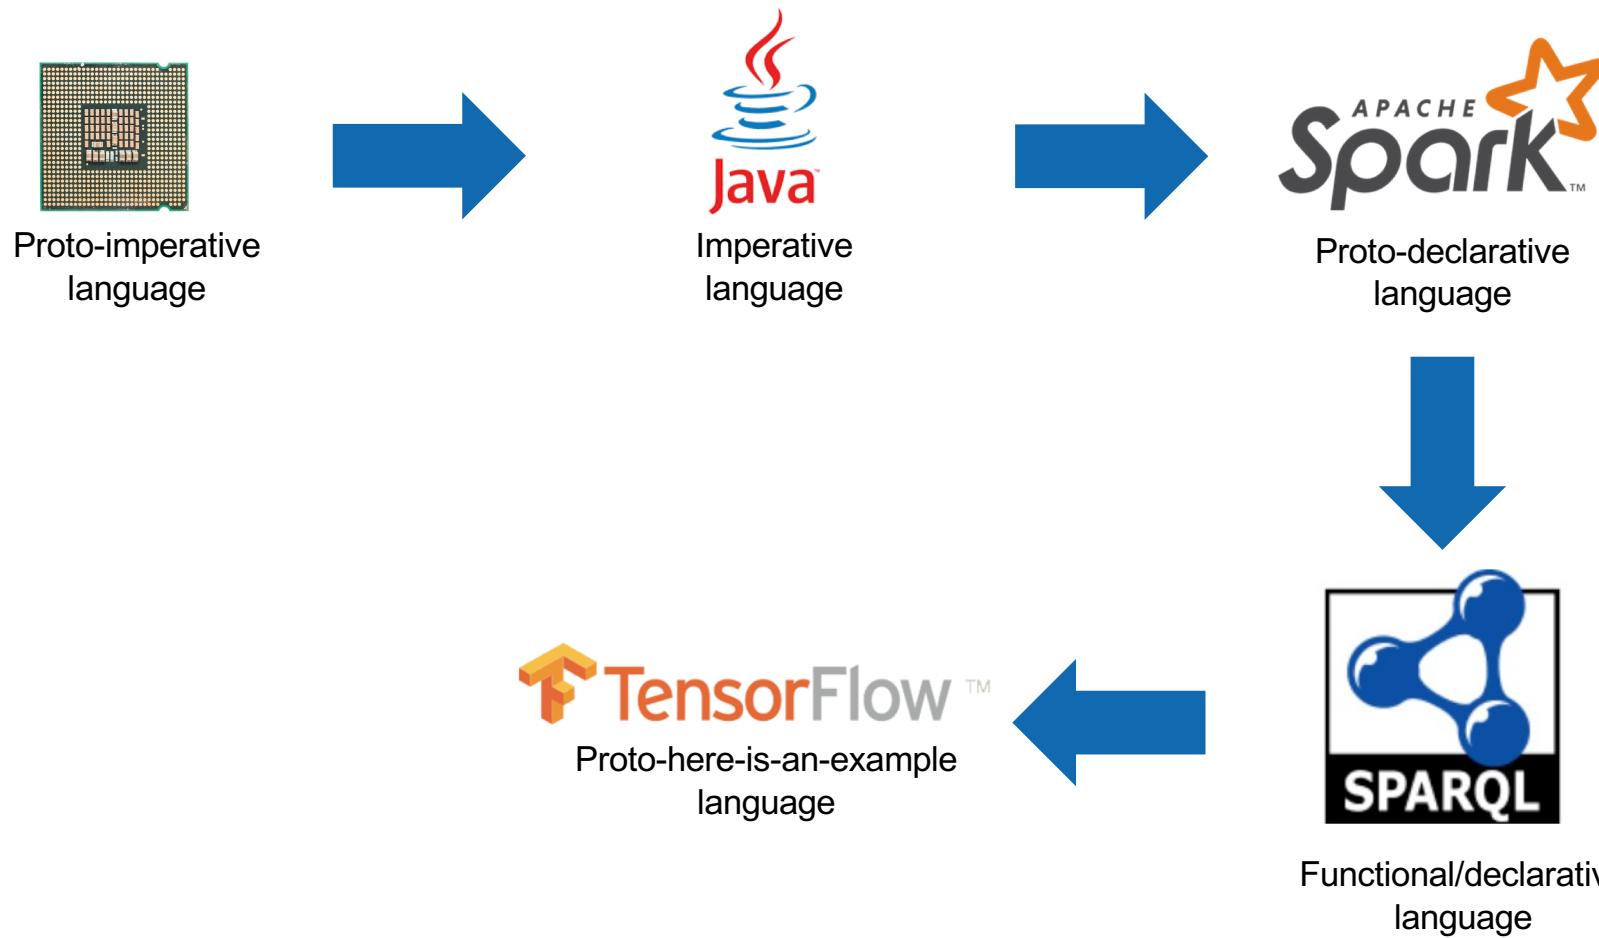


Functional/declarative
language

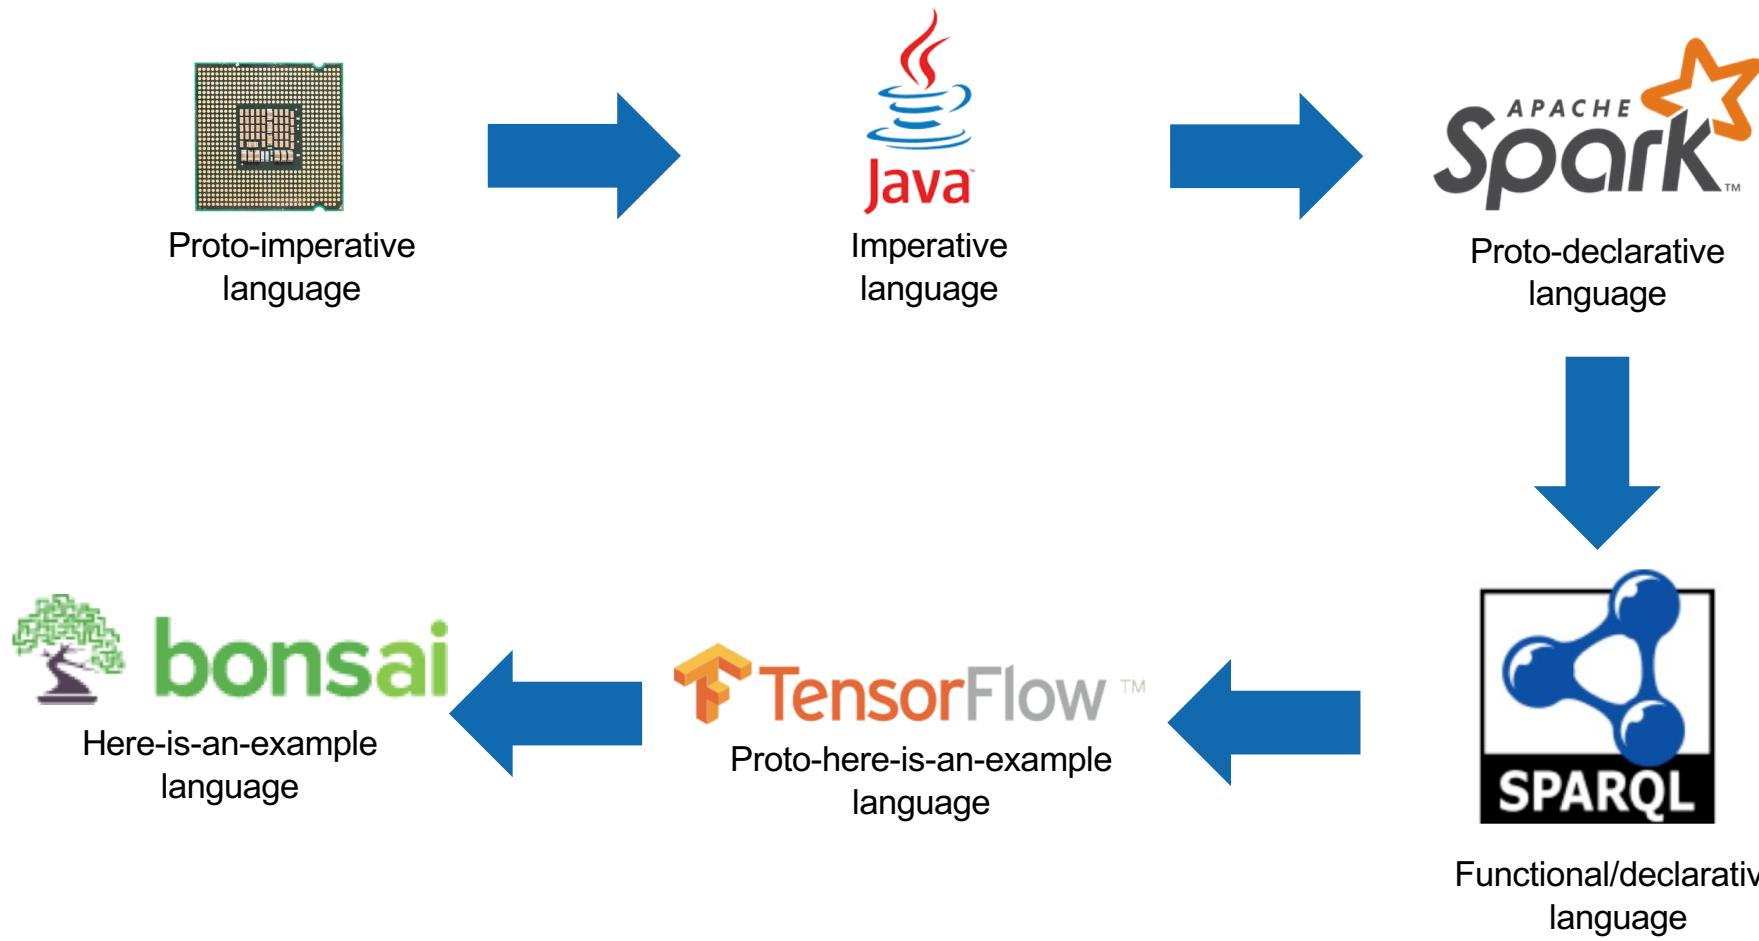
Language landscape



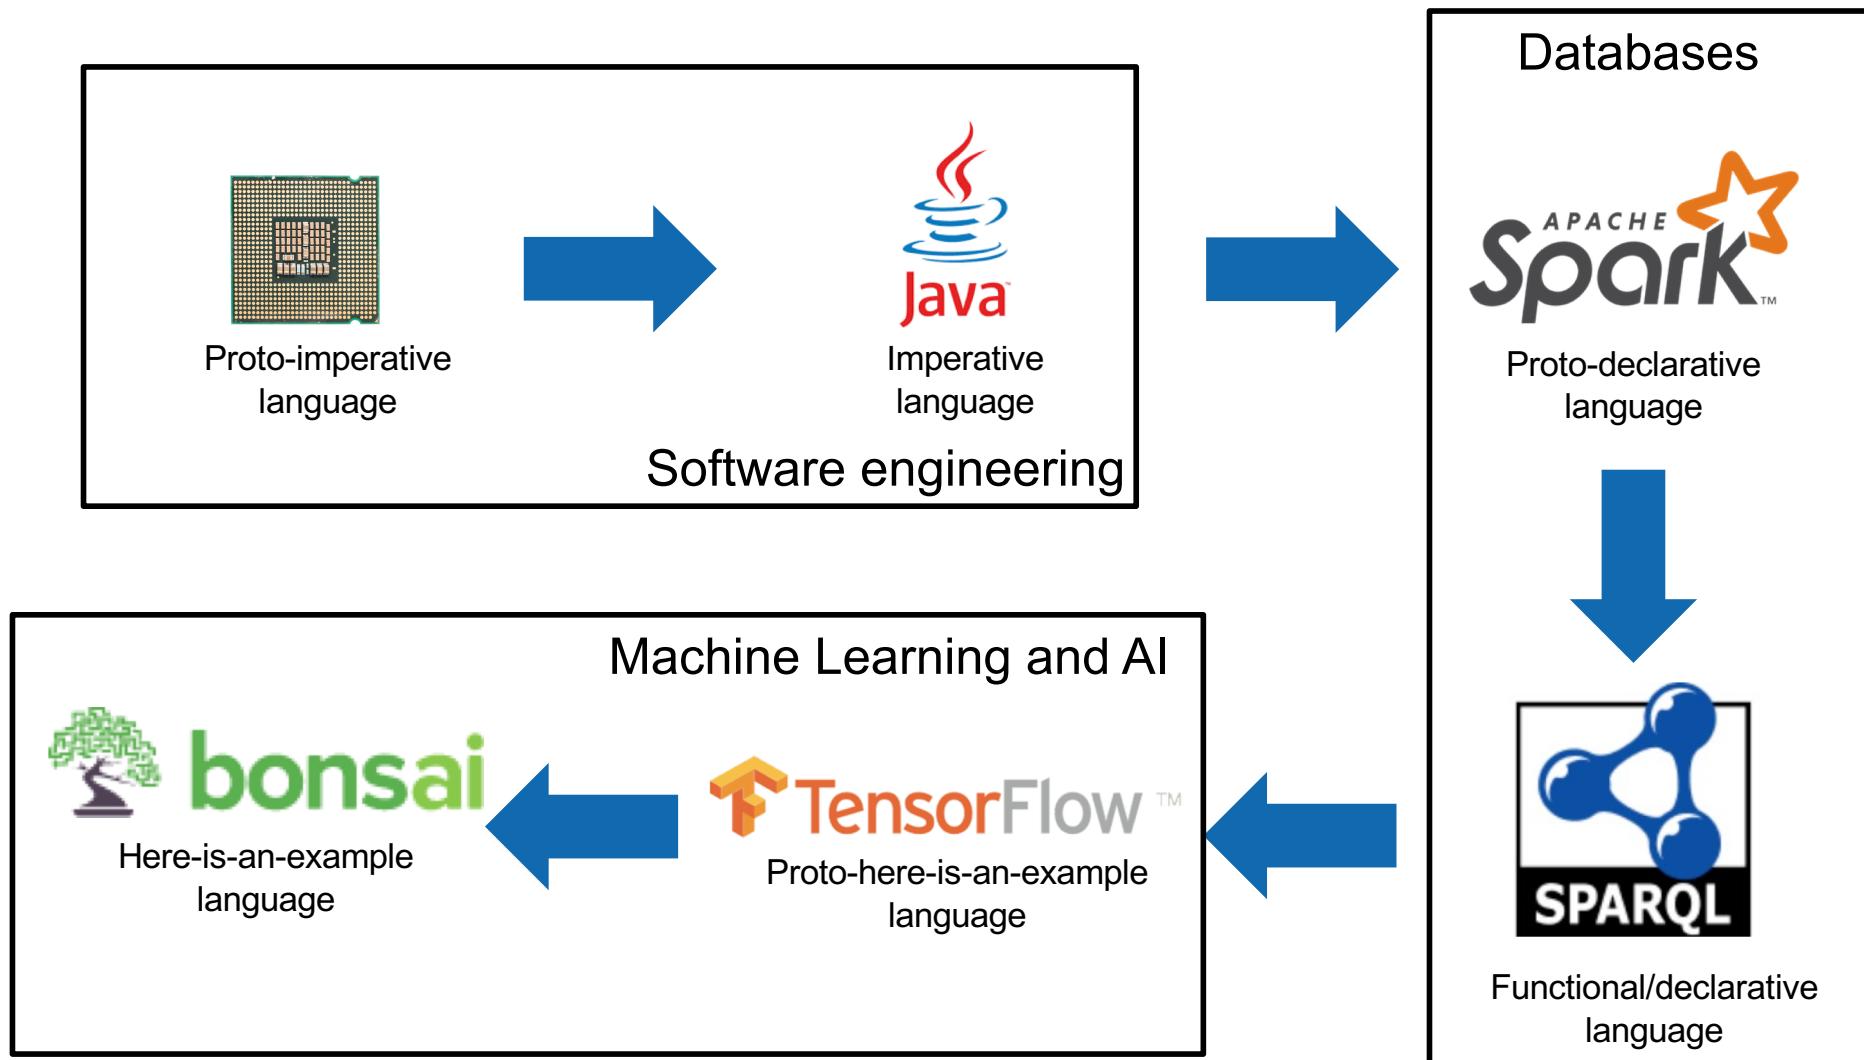
Language landscape



Language landscape



Language landscape



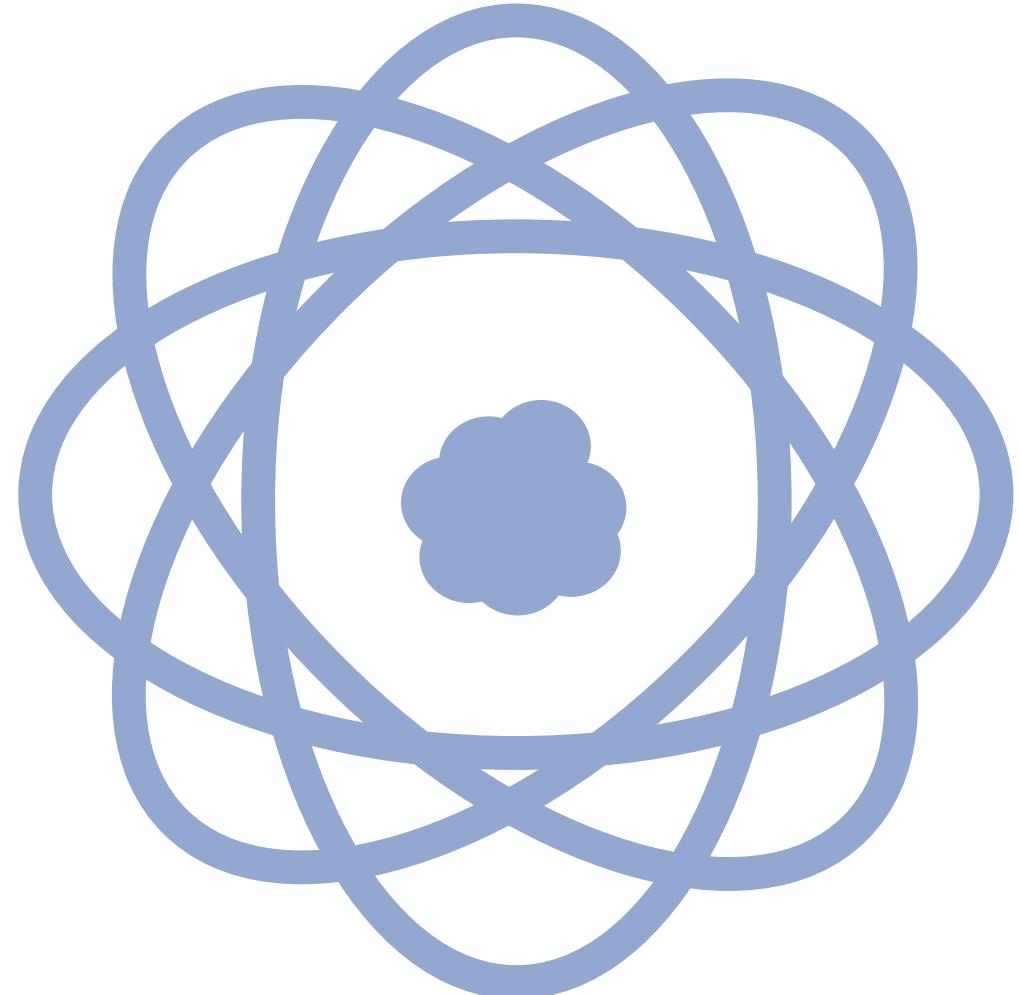


Transactions

The good old times of databases: ACID

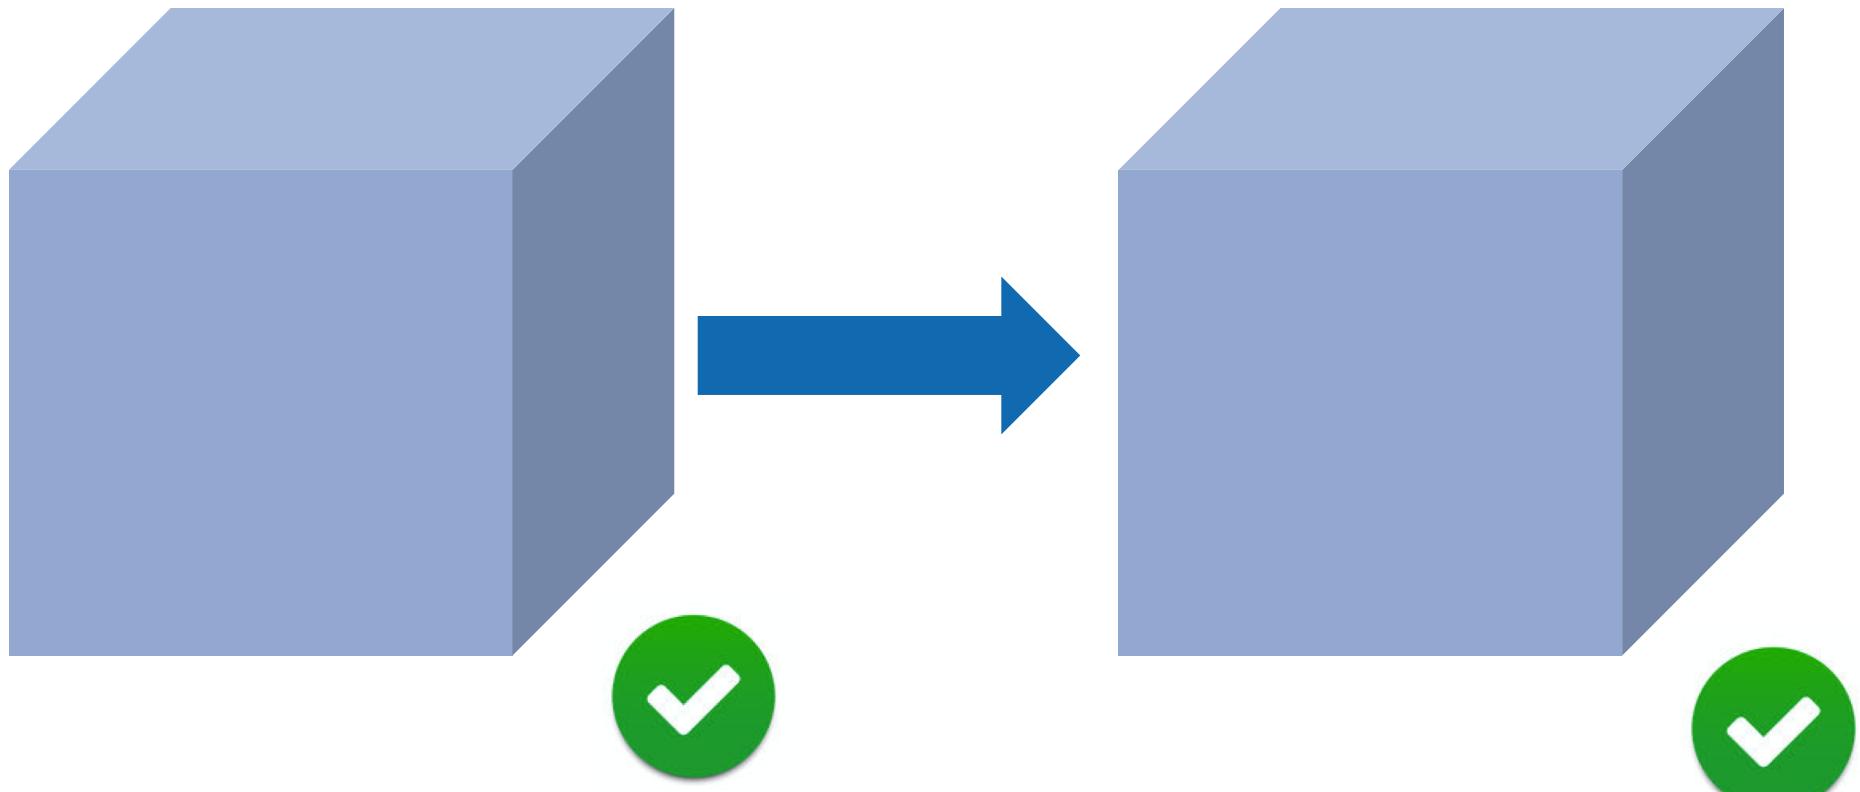
Atomicity
Consistency
Isolation
Durability

Atomicity



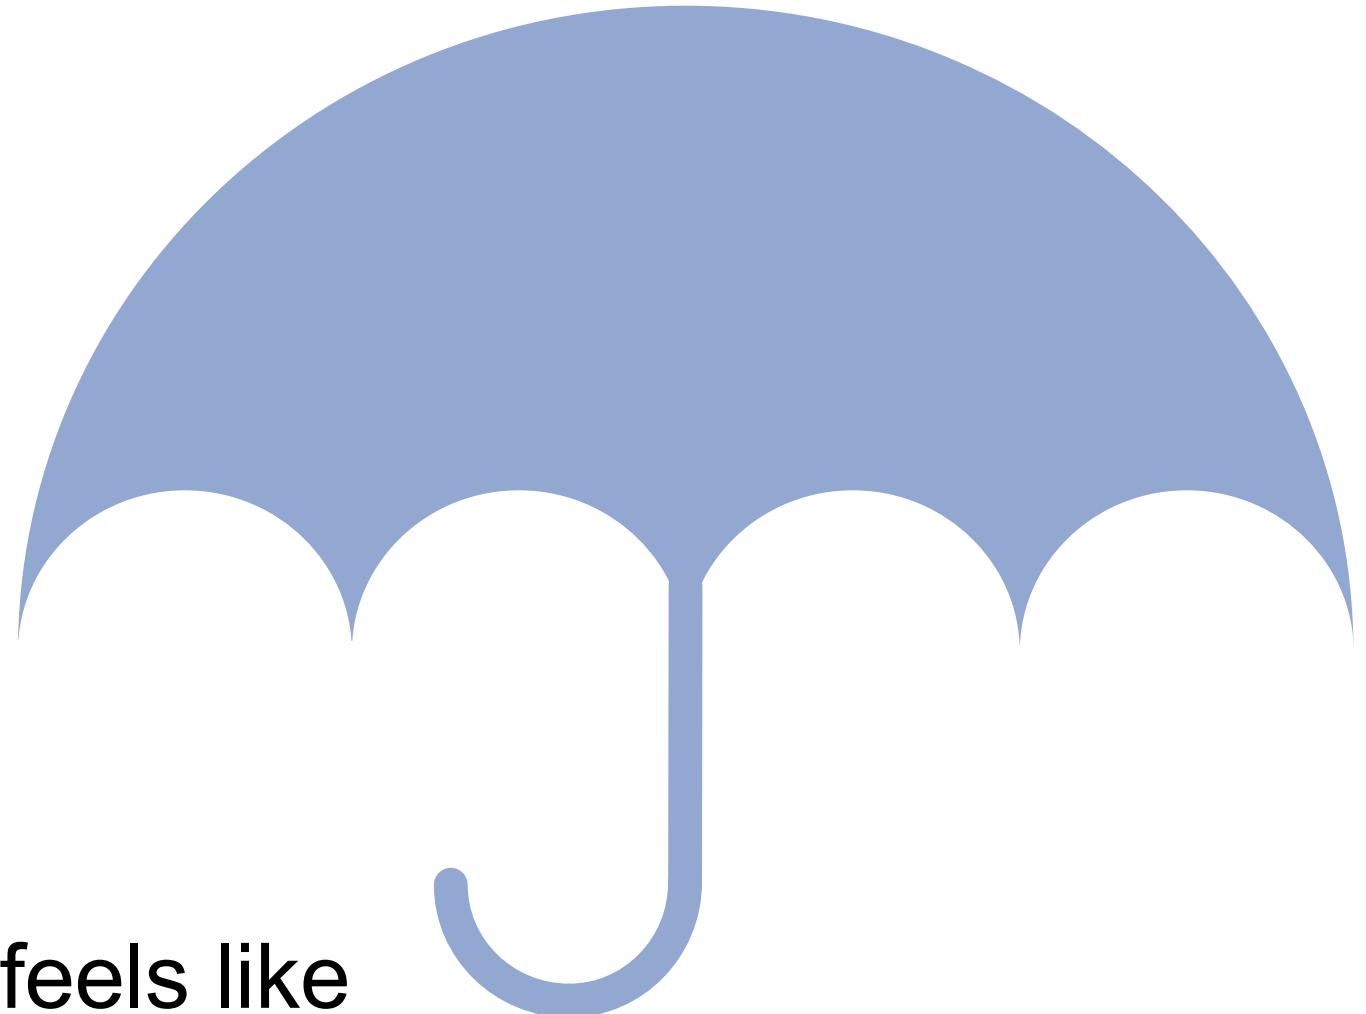
Either the entire transaction is applied,
or none of it (rollback).

Consistency



After a transaction, the database is in a consistent state again.

Isolation



A transaction feels like
nobody else is writing to the database.

Durability

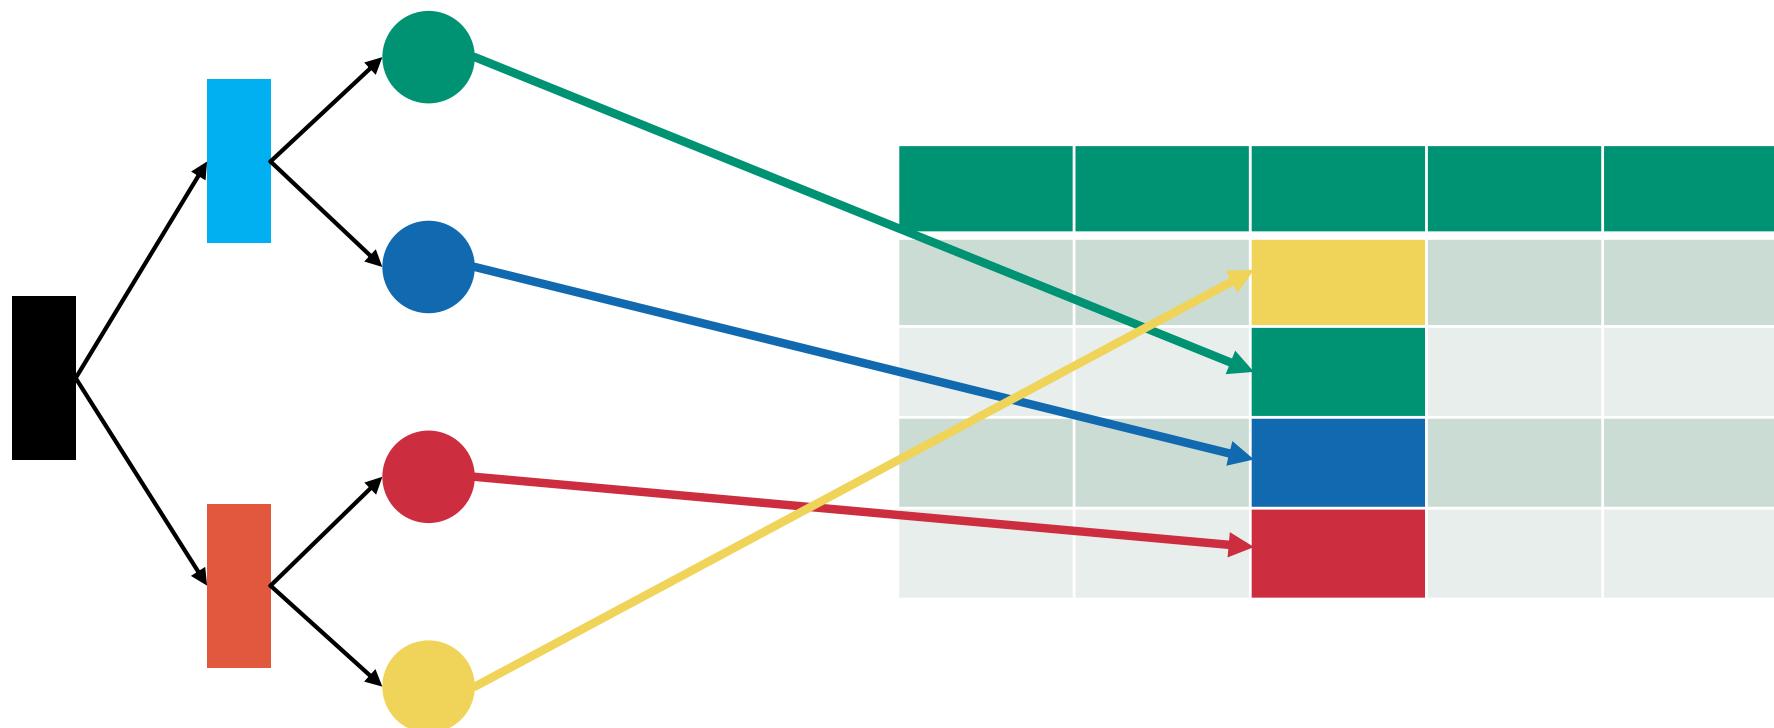


Updates made do not disappear again.



Performance

Indices



Optimize for read vs. write intensive

OnLine
Transaction
Processing

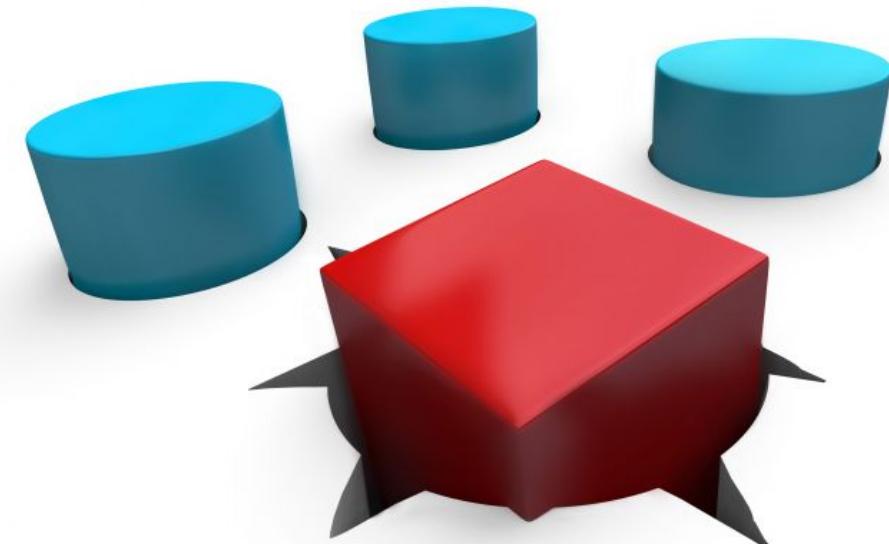


Write-intensive

Read-intensive

No such thing as "one size fits all"

Data shapes
matter





Data Scale-Up

Data can have...

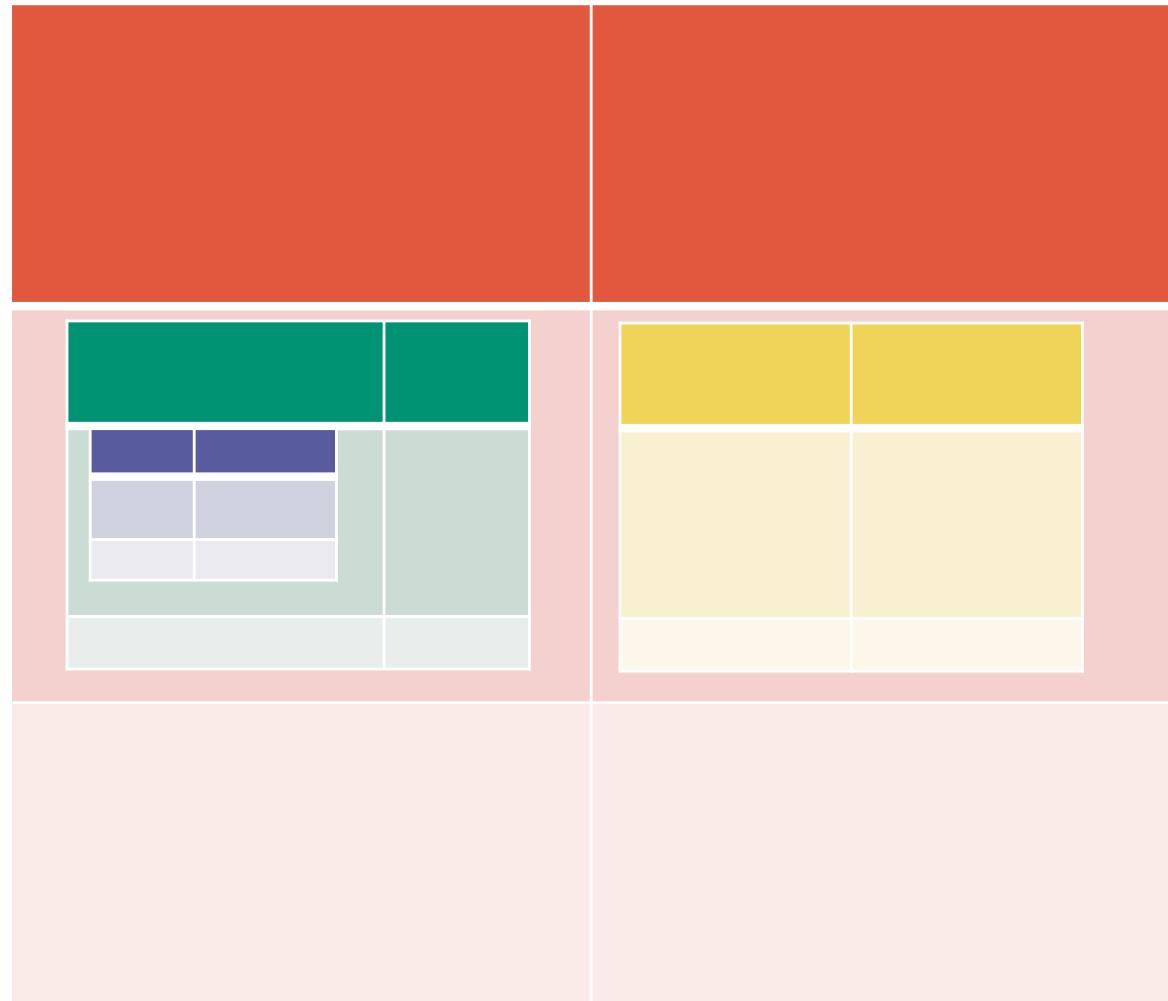
Lots of rows

Data can have...

Lots of columns

Data can have...

Lots of nesting



The rest of the lecture: Scaling up

Lots of rows	Object Storage
Lots of rows	Distributed File Systems
Lots of nesting	Syntax
Lots of rows/columns	Column storage
Lots of nesting	Data Models
Lots of rows	Massive Parallel Processing
Lots of nesting	Document Stores
Lots of nesting	Querying